

2.10 — Introduction to the preprocessor

BY ALEX ON JUNE 3RD, 2007 | LAST MODIFIED BY NASCARDRIVER ON FEBRUARY 4TH, 2020

Translation and the preprocessor

When you compile your code, you might expect that the compiler compiles the code exactly as you've written it. This actually isn't the case.

Prior to compilation, the code file goes through a phase known as **translation**. Many things happen in the translation phase to get your code ready to be compiled (if you're curious, you can find a list of translation phases [here](#)). A code file with translations applied to it is called a **translation unit**.

The most noteworthy of the translation phases involves the preprocessor. The **preprocessor** is best thought of as a separate program that manipulates the text in each code file.

When the preprocessor runs, it scans through the code file (from top to bottom), looking for preprocessor directives. **Preprocessor directives** (often just called *directives*) are instructions that start with a `#` symbol and end with a newline (NOT a semicolon). These directives tell the preprocessor to perform specific particular text manipulation tasks. Note that the preprocessor does not understand C++ syntax -- instead, the directives have their own syntax (which in some cases resembles C++ syntax, and in other cases, not so much).

The output of the preprocessor goes through several more translation phases, and then is compiled. Note that the preprocessor does not modify the original code files in any way -- rather, all text changes made by the preprocessor happen temporarily in-memory each time the code file is compiled.

In this lesson, we'll discuss what some of the most common preprocessor directives do.

As an aside...

Using directives (introduced in lesson [2.9 -- Naming collisions and an introduction to namespaces](#)) are not preprocessor directives (and thus are not processed by the preprocessor). So while the term *directive* *usually* means a preprocessor directive, this is not always the case.

Includes

You've already seen the `#include` directive in action (generally to `#include <iostream>`). When you `#include` a file, the preprocessor replaces the `#include` directive with the contents of the included file. The included contents are then preprocessed (along with the rest of the file), and then compiled.

Consider the following program:

```
1  #include <iostream>
2
3  int main()
4  {
5      std::cout << "Hello, world!";
6      return 0;
7  }
```

When the preprocessor runs on this program, the preprocessor will replace `#include <iostream>` with the preprocessed contents of the file named "iostream".

Since `#include` is almost exclusively used to include header files, we'll discuss `#include` in more detail in the next lesson (when we discuss header files in more detail).

Macro defines

The `#define` directive can be used to create a macro. In C++, a **macro** is a rule that defines how input text is converted into replacement output text.

There are two basic types of macros: *object-like macros*, and *function-like macros*.

Function-like macros act like functions, and serve a similar purpose. We will not discuss them here, because their use is generally considered dangerous, and almost anything they can do can be done by a normal function.

Object-like macros can be defined in one of two ways:

```
#define identifier
#define identifier substitution_text
```

The top definition has no substitution text, whereas the bottom one does. Because these are preprocessor directives (not statements), note that neither form ends with a semicolon.

Object-like macros with substitution text

When the preprocessor encounters this directive, any further occurrence of the identifier is replaced by *substitution_text*. The identifier is traditionally typed in all capital letters, using underscores to represent spaces.

Consider the following program:

```
1  #include <iostream>
2
3  #define MY_NAME "Alex"
4
5  int main()
6  {
7      std::cout << "My name is: " << MY_NAME;
8
9      return 0;
10 }
```

The preprocessor converts the above into the following:

```
1  // The contents of iostream are inserted here
2
3  int main()
4  {
5      std::cout << "My name is: " << "Alex";
6
7      return 0;
8  }
```

Which, when run, prints the output `My name is: Alex`.

We recommend avoiding these kinds of macros altogether, as there are better ways to do this kind of thing. We discuss this more in lesson [4.13 -- Const, constexpr, and symbolic constants](#).

Object-like macros without substitution text

Object-like macros can also be defined without substitution text.

For example:

```
1 | #define USE_YEN
```

Macros of this form work like you might expect: any further occurrence of the identifier is removed and replaced by nothing!

This might seem pretty useless, and it *is* *useless* for doing text substitution. However, that's not what this form of the directive is generally used for. We'll discuss the uses of this form in just a moment.

Unlike object-like macros with substitution text, macros of this form are generally considered acceptable to use.

Conditional compilation

The *conditional compilation* preprocessor directives allow you to specify under what conditions something will or won't compile. There are quite a few different conditional compilation directives, but we'll only cover the three that are used by far the most here: `#ifdef`, `#ifndef`, and `#endif`.

The `#ifdef` preprocessor directive allows the preprocessor to check whether an identifier has been previously `#defined`. If so, the code between the `#ifdef` and matching `#endif` is compiled. If not, the code is ignored.

Consider the following program:

```
1 | #include <iostream>
2 |
3 | #define PRINT_JOE
4 |
5 | int main()
6 | {
7 |     #ifdef PRINT_JOE
8 |         std::cout << "Joe\n"; // if PRINT_JOE is defined, compile this code
9 |     #endif
10 |
11 |     #ifdef PRINT_BOB
12 |         std::cout << "Bob\n"; // if PRINT_BOB is defined, compile this code
13 |     #endif
14 |
15 |     return 0;
16 | }
```

Because `PRINT_JOE` has been `#defined`, the line `cout << "Joe\n"` will be compiled. Because `PRINT_BOB` has not been `#defined`, the line `cout << "Bob\n"` will be ignored.

`#ifndef` is the opposite of `#ifdef`, in that it allows you to check whether an identifier has *NOT* been `#defined` yet.

```
1 | #include <iostream>
2 |
3 | int main()
4 | {
5 |     #ifndef PRINT_BOB
6 |         std::cout << "Bob\n";
7 |     #endif
8 |
9 |     return 0;
10 | }
```

This program prints "Bob", because `PRINT_BOB` was never `#defined`.

#if 0

One more common use of conditional compilation involves using `#if 0` to exclude a block of code from being compiled (as if it were inside a comment block):

```
1  #include <iostream>
2
3  int main()
4  {
5      std::cout << "Joe\n";
6
7      #if 0 // Don't compile anything starting here
8          std::cout << "Bob\n";
9          std::cout << "Steve\n";
10     #endif // until this point
11
12     return 0;
13 }
```

The above code only prints “Joe”, because “Bob” and “Steve” were inside an `#if 0` block that the preprocessor will exclude from compilation.

This provides a convenient way to “comment out” code that contains multi-line comments.

Object-like macros don't affect other preprocessor directives

Now you might be wondering:

```
1  #define PRINT_JOE
2
3  #ifdef PRINT_JOE
4  // ...
```

Since we defined `PRINT_JOE` to be nothing, how come the preprocessor didn't replace `PRINT_JOE` in `#ifdef PRINT_JOE` with nothing?

Macros only cause text substitution for normal code. Other preprocessor commands are ignored. Consequently, the `PRINT_JOE` in `#ifdef PRINT_JOE` is left alone.

For example:

```
1  #define F00 9 // Here's a macro substitution
2
3  #ifdef F00 // This F00 does not get replaced because it's part of another preprocessor directive
4      ve
5      std::cout << F00; // This F00 gets replaced with 9 because it's part of the normal code
6  #endif
```

In actuality, the output of the preprocessor contains no directives at all -- they are all resolved/stripped out before compilation, because the compiler wouldn't know what to do with them.

The scope of defines

Directives are resolved before compilation, from top to bottom on a file-by-file basis.

Consider the following program:

```
1  #include <iostream>
2
```

```
3 void foo()
4 {
5     #define MY_NAME "Alex"
6 }
7
8 int main()
9 {
10     std::cout << "My name is: " << MY_NAME;
11
12     return 0;
13 }
```

Even though it looks like `#define MY_NAME "Alex"` is defined inside function `foo`, the preprocessor won't notice, as it doesn't understand C++ concepts like functions. Therefore, this program behaves identically to one where `#define MY_NAME "Alex"` was defined either before or immediately after function `foo`. For general readability, you'll generally want to `#define` identifiers outside of functions.

Once the preprocessor has finished, all defined identifiers from that file are discarded. This means that directives are only valid from the point of definition to the end of the file in which they are defined. Directives defined in one code file do not have impact on other code files in the same project.

Consider the following example:

function.cpp:

```
1 #include <iostream>
2
3 void doSomething()
4 {
5     #ifdef PRINT
6         std::cout << "Printing!";
7     #endif
8     #ifndef PRINT
9         std::cout << "Not printing!";
10    #endif
11 }
```

main.cpp:

```
1 void doSomething(); // forward declaration for function doSomething()
2
3 #define PRINT
4
5 int main()
6 {
7     doSomething();
8
9     return 0;
10 }
```

The above program will print:

Not printing!

Even though `PRINT` was defined in `main.cpp`, that doesn't have any impact on any of the code in `function.cpp` (`PRINT` is only `#defined` from the point of definition to the end of `main.cpp`). This will be of consequence when we discuss header guards in a future lesson.

[2.11 -- Header files](#)[Index](#)[2.9 -- Naming collisions and an introduction to namespaces](#)[C++ TUTORIAL](#) | [PRINT THIS POST](#)

217 comments to 2.10 — Introduction to the preprocessor

[« Older Comments](#) [1](#) [2](#) [3](#) [4](#)

Uros

[February 5, 2020 at 1:38 am · Reply](#)

Why does this program print the ifNdef text instead of ifdef text?

```
1  #include <iostream>
2  void Zdravo()
3  {
4      #ifdef HELLO
5          std::cout << "Zdravo tamo!";
6      #endif
7  }
8  void NeBasZdravo()
9  {
10     #ifndef HELLO
11         std::cout << "Ne bas tako zdravo tamo!";
12     #endif
13 }
14 int main()
15 {
16     #define HELLO
17     Zdravo();
18     NeBasZdravo();
19     return 0;
20 }
```



nascardriver

[February 6, 2020 at 9:09 am · Reply](#)

The preprocessor doesn't care about control flow, it processes files from top to bottom. The #define in line 16 doesn't affect the #ifs.



Uros

[February 7, 2020 at 3:44 am · Reply](#)

what does that mean exactly? what should i do to fix the problem?



nascar driver

[February 7, 2020 at 8:27 am · Reply](#)

```

1  #include <iostream>
2
3  // Move the define to a place before its first use
4  #define HELLO
5
6  void Zdravo()
7  {
8      #ifdef HELLO
9      std::cout << "Zdravo tamo!";
10     #endif
11 }
12
13 void NeBasZdravo()
14 {
15     #ifndef HELLO
16     std::cout << "Ne bas tako zdravo tamo!";
17     #endif
18 }
19
20 int main()
21 {
22     Zdravo();
23     NeBasZdravo();
24     return 0;
25 }
```



Uros

[February 8, 2020 at 1:05 am · Reply](#)

oh, i understand, well i think i do at least, thanks bro



Ryan

[February 2, 2020 at 9:33 pm · Reply](#)

Quick typo! In the 'Conditional compilation' section- "The #ifdef preprocessor directive allow the preprocessor"; 'allow' instead of 'allows'.



nascar driver

[February 4, 2020 at 8:26 am · Reply](#)

Fixed, thanks!



Apaulture

[January 15, 2020 at 2:44 pm · Reply](#)

In regards to conditional compilation preprocessor directives,

```

1  #ifdef MACRO
2      // statement
3  #endif
```

is useful for saving the entire process some overhead in the case that a MACRO is not defined since we wouldn't have to resort to phase 5-9 and eventually the compilation process.

```
1 | #ifndef MACRO
2 |     // statement
3 | #endif
```

may be useful as a backup tool (to run the needed code) in the case a MACRO is not defined or missing.



aqib

December 17, 2019 at 3:56 am · Reply

I could not understand this topic properly
i just want to know if i do not remember things of this topic so can i use c++ without these things



nascar driver

December 17, 2019 at 5:14 am · Reply

You can continue, but you'll have to come back eventually.



Omran

January 27, 2020 at 3:54 am · Reply

hey i can help maybe , i watched this video on youtube , <https://www.youtube.com/watch?v=j3mYki1SrKE> , it explains part of this lesson :)



koe

December 9, 2019 at 3:47 pm · Reply


The previous section on namespaces mentioned 'using' is a directive. Is this accurate, and how does the preprocessor handle 'using' instances?



nascar driver

December 11, 2019 at 4:31 am · Reply

They're indeed called using-directives, but they're not handled by the preprocessor.




Chayim

November 25, 2019 at 12:26 am · Reply

Why does it print 'Not printing' ?

If it's in the same file function.cpp as 'Printing' that is not being compiled because it was not defined because define in main.pp does not have affect on function.pp file, so is 'Not printing' in this same file as 'Printing'.




Alex

November 25, 2019 at 10:22 pm · Reply

PRINT is only #defined from the point of definition to the end of main.cpp.

When function.cpp is being compiled, PRINT has not been #defined as far as function.cpp is concerned, so the #ifndef directive compiles.



Chayim

[November 25, 2019 at 10:47 pm · Reply](#)

My mistake was that I thought that the 'Not printing' is directed as `#ifdef` like 'Printing' but now I recognized that it's `#ifndef PRINT`.



hellmet

[November 6, 2019 at 11:30 am · Reply](#)

I was just reviewing...to make sure I retained everything :)

In the 'conditional compilation' section's example, the comment reads 'execute this code', which is somewhat misleading. Later on, the tutorial correctly points out the code is not compiled at all. I think the 'conditional compilation' example should be amended to show this fact.



nascar driver

[November 7, 2019 at 2:46 am · Reply](#)

The "execute this code" was connected to the comment in the line above. I moved the entire comment into one line, that should be easier to understand. Thanks for the suggestion :)



hellmet

[November 7, 2019 at 4:09 am · Reply](#)

Ohh right, that makes more sense!

I meant, instead of 'execute this code', 'compile this code' would reflect the preprocessor nature of the example. But yes, now it's more clearer.

My pleasure!



nascar driver

[November 7, 2019 at 4:13 am · Reply](#)

Agreed and updated!



hellmet

[November 7, 2019 at 4:14 am · Reply](#)

My pleasure and glad to be of help! :)

[« Older Comments](#)

[1](#) [2](#) [3](#) [4](#)