

5.7 — For statements

BY ALEX ON JUNE 25TH, 2007 | LAST MODIFIED BY ALEX ON JANUARY 23RD, 2020

By far, the most utilized looping statement in C++ is the *for statement*. The **for statement** (also called a **for loop**) is ideal when we know exactly how many times we need to iterate, because it lets us easily define, initialize, and change the value of loop variables after each iteration.

The *for statement* looks pretty simple in abstract:

```
for (init-statement; condition-expression; end-expression)
    statement
```

The easiest way to understand a *for loop* is to convert it into an equivalent *while loop*:

```
{ // note the block here
    init-statement;
    while (condition-expression)
    {
        statement;
        end-expression;
    }
} // variables defined inside the loop go out of scope here
```

The variables defined inside a *for loop* have a special kind of scope called *loop scope*. Variables with **loop scope** exist only within the loop, and are not accessible outside of it.

Evaluation of for statements

A *for statement* is evaluated in 3 parts:

- 1) The init-statement is evaluated. Typically, the init-statement consists of variable definitions and initialization. This statement is only evaluated once, when the loop is first executed.
- 2) The condition-expression is evaluated. If this evaluates to false, the loop terminates immediately. If this evaluates to true, the statement is executed.
- 3) After the statement is executed, the end-expression is evaluated. Typically, this expression is used to increment or decrement the variables declared in the init-statement. After the end-expression has been evaluated, the loop returns to step 2.

Let's take a look at a sample *for loop* and discuss how it works:

```
1 | for (int count{ 0 }; count < 10; ++count)
2 |     std::cout << count << ' ';
```

First, we declare a loop variable named `count`, and assign it the value 0.

Second, `count < 10` is evaluated, and since `count` is 0, `0 < 10` evaluates to true. Consequently, the statement executes, which prints 0.

Third, `++count` is evaluated, which increments `count` to 1. Then the loop goes back to the second step.

Now, `1 < 10` is evaluated to true, so the loop iterates again. The statement prints 1, and `count` is incremented to 2.

`2 < 10` evaluates to true, the statement prints 2, and count is incremented to 3. And so on.

Eventually, count is incremented to 10, `10 < 10` evaluates to false, and the loop exits.

Consequently, this program prints the result:

0 1 2 3 4 5 6 7 8 9

For loops can be hard for new programmers to read -- however, experienced programmers love them because they are a very compact way to do loops of this nature. For the sake of example, let's uncompact the above *for loop* by converting it into an equivalent *while loop*:

```
1 { // outer braces ensure loop scope
2   int count{ 0 };
3   while (count < 10)
4   {
5       std::cout << count << ' ';
6       ++count;
7   }
8 }
```

That doesn't look so bad, does it? Note that the outer braces are necessary here, because count goes out of scope when the loop ends.

More for loop examples

Here's an example of a *for loop* used to calculate an exponentiation of integers:

```
1 // returns the value nBase ^ nExp
2 int pow(int base, int exponent)
3 {
4     int total{ 1 };
5
6     for (int count{ 0 }; count < exponent; ++count)
7         total *= base;
8
9     return total;
10 }
```

This function returns the value $\text{base}^{\text{exponent}}$ (base to the exponent power).

This is a straightforward incrementing *for loop*, with count looping from 0 up to (but excluding) exponent.

If exponent is 0, the *for loop* will execute 0 times, and the function will return 1.

If exponent is 1, the *for loop* will execute 1 time, and the function will return $1 * \text{base}$.

If exponent is 2, the *for loop* will execute 2 times, and the function will return $1 * \text{base} * \text{base}$.

Although most *for loops* increment the loop variable by 1, we can decrement it as well:

```
1 for (int count{ 9 }; count >= 0; --count)
2     std::cout << count << ' ';
```

This prints the result:

9 8 7 6 5 4 3 2 1 0

Alternately, we can change the value of our loop variable by more than 1 with each iteration:

```
1 for (int count{ 9 }; count >= 0; count -= 2)
2     std::cout << count << ' ';
```

This prints the result:

9 7 5 3 1

Off-by-one errors

One of the biggest problems that new programmers have with *for loops* (and other kinds of loops) is off-by-one errors. **Off-by-one errors** occur when the loop iterates one too many or one too few times. This generally happens because the wrong relational operator is used in the conditional-expression (eg. > instead of >=). These errors can be hard to track down because the compiler will not complain about them -- the program will run fine, but it will produce the wrong result.

When writing *for loops*, remember that the loop will execute as long as the conditional-expression is true. Generally it is a good idea to test your loops using known values to make sure that they work as expected. A good way to do this is to test your loop with known inputs that cause it to iterate 0, 1, and 2 times. If it works for those, it will likely work for any number of iterations.

Rule: Test your loops with known inputs that cause it to iterate 0, 1, and 2 times.

Omitted expressions

It is possible to write *for loops* that omit any or all of the expressions. For example, in the following example, we'll omit the init-statement and end-expression:

```
1  int count=0;
2  for ( ; count < 10; )
3  {
4      std::cout << count << ' ';
5      ++count;
6  }
```

This *for loop* produces the result:

0 1 2 3 4 5 6 7 8 9

Rather than having the *for loop* do the initialization and incrementing, we've done it manually. We have done so purely for academic purposes in this example, but there are cases where not declaring a loop variable (because you already have one) or not incrementing it (because you're incrementing it some other way) are desired.

Although you do not see it very often, it is worth noting that the following example produces an infinite loop:

```
for (;;)
    statement;
```

The above example is equivalent to:

```
while (true)
    statement;
```

This might be a little unexpected, as you'd probably expect an omitted condition-expression to be treated as "false". However, the C++ standard explicitly (and inconsistently) defines that an omitted condition-expression in a *for loop* should be treated as "true".

We recommend avoiding this form of the *for loop* altogether and using *while(true)* instead.

Multiple declarations

Although for loops typically iterate over only one variable, sometimes *for loops* need to work with multiple variables. When this happens, the programmer can make use of the comma operator in order to assign (in the init-statement) or change (in the end-statement) the value of multiple variables:

```
1 int iii{};
2 int jjj{};
3 for (iii = 0, jjj = 9; iii < 10; ++iii, --jjj)
4     std::cout << iii << ' ' << jjj << '\n';
```

This loop assigns values to two previously declared variables: `iii` to 0, and `jjj` to 9. It iterates `iii` over the range 0 to 9, and each iteration `iii` is incremented and `jjj` is decremented.

This program produces the result:

```
0 9
1 8
2 7
3 6
4 5
5 4
6 3
7 2
8 1
9 0
```

Note: More typically, we'd write the above loop as:

```
1 for (int iii{ 0 }, jjj{ 9 }; iii < 10; ++iii, --jjj)
2     std::cout << iii << ' ' << jjj << '\n';
```

In this case, the comma in the init-statement is part of the variable definition syntax, not a use of the comma operator. But the effect is identical.

For loops in old code

In older versions of C++, variables defined as part of the init-statement did not get destroyed at the end of the loop. This meant that you could have something like this:

```
1 for (int count{ 0 }; count < 10; ++count) // count defined here
2     std::cout << count << ' ';
3
4 // count is not destroyed in older compilers
5
6 std::cout << '\n';
7 std::cout << "I counted to: " << count << '\n'; // so you can still use it here
```

This use has been disallowed, but you may still see it in older code.

Nested for loops

Like other types of loops, for loops can be nested inside other loops. In the following example, we're nesting a for loop inside another for loop:

```
1 #include <iostream>
2
3 int main()
4 {
5     for (char c{ 'a' }; c <= 'e'; ++c) // outer loop on letters
6     {
7         std::cout << c; // print our letter first
```

```

8
9         for (int i{ 0 }; i < 3; ++i) // inner loop on all numbers
10            std::cout << i;
11
12            std::cout << '\n';
13        }
14
15        return 0;
16    }

```

For each iteration of the outer loop, the inner loop runs in its entirety. Consequently, the output is:

```

a012
b012
c012
d012
e012

```

Here's some more detail on what's happening here. The outer loop runs first, and char `c` is initialized to 'a'. Then `c <= 'e'` is evaluated, which is true, so the loop body executes. Since `c` is set to 'a', this first prints 'a'. Next the inner loop executes entirely (which prints '0', '1', and '2'). Then a newline is printed. Now the outer loop body is finished, so the outer loop returns to the top, `c` is incremented to 'b', and the loop condition is re-evaluated. Since the loop condition is still true the next iteration of the outer loop begins. This prints ("b012\n"). And so on.

Conclusion

For statements are the most commonly used loop in the C++ language. Even though its syntax is typically a bit confusing to new programmers, you will see *for loops* so often that you will understand them in no time at all!

Quiz

- 1) Write a *for loop* that prints every even number from 0 to 20.
- 2) Write a function named `sumTo()` that takes an integer parameter named `value`, and returns the sum of all the numbers from 1 to `value`.

For example, `sumTo(5)` should return 15, which is $1 + 2 + 3 + 4 + 5$.

Hint: Use a non-loop variable to accumulate the sum as you iterate from 1 to the input value, much like the `pow()` example above uses the `total` variable to accumulate the return value each iteration.

- 3) What's wrong with the following *for loop*?

```

1 // Print all numbers from 9 to 0
2 for (unsigned int count{ 9 }; count >= 0; --count)
3     std::cout << count << ' ';

```

Quiz solutions

- 1) [Show Solution](#)
- 2) [Show Solution](#)
- 3) [Show Solution](#)



5.8 -- Break and continue

[Index](#)[5.6 -- Do while statements](#)[C++ TUTORIAL](#) | [PRINT THIS POST](#)

229 comments to 5.7 — For statements

[« Older Comments](#)[1](#) [2](#) [3](#) [4](#)**Eric**[January 16, 2020 at 7:47 am · Reply](#)

Here is my code for Quiz question 2).

Is it considered a good or bad idea to name the input/variables from one function to another with the same name when they reference the same thing?

```
#include <iostream>
```

```
int getNumber()
```

```
{
    std::cout << "Input an integer you wish to sum to: ";
    int value{};
    std::cin >> value;
    return value;
}
```

```
int calcSum(int value)
```

```
{
    int total{};
    for (int count{ 1 }; count <= value; ++count)
        total += count;
    return total;
}
```

```
void printTotal (int value , int total)
```

```
{
    std::cout << "The sum of the number from 1 to " << value << " is " << total << "." << '\n';
}
```

```
int main()
```

```
{
    int value{ getNumber() };
    int total{ calcSum( value ) };

    printTotal( value , total );

    return 0;
}
```

nascardriver



January 16, 2020 at 7:51 am · Reply

Please use code tags when posting code.

If two variables mean the same thing, they should have the same name. There are exceptions when a variable's name is affected by its context.

```

1  #include <iostream>
2  #include <string>
3
4  // @strTitle is obviously the title of the book
5  // "strBookTitle" would be ambiguous
6  void printBook(std::string strTitle)
7  {
8      std::cout << strTitle << '\n';
9  }
10
11 int main()
12 {
13     // Here we don't know what the title would be for
14     // if we named the variables "strTitle", so we add
15     // "book".
16     std::string strBookTitle{ "Eric's adventures" };
17
18     printBook(strBookTitle);
19
20     return 0;
21 }
```



kavin

January 16, 2020 at 2:26 am · Reply

Under "multiple declarations" my compiler keep showing this error if i use uniform initialization like in the example.

```

1  for (iii{ 0 }, jjj{ 9 }; iii < 10; ++iii, --jjj) //shows error
```

error C2064: term does not evaluate to a function taking 1 arguments

error C2064: term does not evaluate to a function taking 1 arguments

I use VS2019. It compiles only if i use uniform or copy initialization like this,

```

1  for (int iii{ 0 }, jjj{ 9 }; iii < 10; ++iii, --jjj)
2      std::cout << iii << ' ' << jjj << '\n';
```

```

1  {
2      int iii, jjj;
3      for (iii=0, jjj=9; iii < 10; ++iii, --jjj) //why only copy init works for me?
4          std::cout << iii << ' ' << jjj << '\n';
5
6      return 0;
7  }
```



nascardriver

January 16, 2020 at 2:36 am · Reply

That was an error in the lesson, thanks for bringing it to our attention!

`iii` and `jjj` already existed, so they can't be initialized anymore. We need to assign a value instead.

[« Older Comments](#)[1](#)[2](#)[3](#)[4](#)