

5.9 — Random number generation

BY ALEX ON DECEMBER 7TH, 2008 | LAST MODIFIED BY ALEX ON JANUARY 23RD, 2020

The ability to generate random numbers can be useful in certain kinds of programs, particularly in games, statistics modeling programs, and scientific simulations that need to model random events. Take games for example -- without random events, monsters would always attack you the same way, you'd always find the same treasure, the dungeon layout would never change, etc... and that would not make for a very good game.

So how do we generate random numbers? In real life, we often generate random results by doing things like flipping a coin, rolling a dice, or shuffling a deck of cards. These events involve so many physical variables (e.g. gravity, friction, air resistance, momentum, etc...) that they become almost impossible to predict or control, and produce results that are for all intents and purposes random.

However, computers aren't designed to take advantage of physical variables -- your computer can't toss a coin, throw a dice, or shuffle real cards. Computers live in a controlled electrical world where everything is binary (false or true) and there is no in-between. By their very nature, computers are designed to produce results that are as predictable as possible. When you tell the computer to calculate $2 + 2$, you *always* want the answer to be 4. Not 3 or 5 on occasion.

Consequently, computers are generally incapable of generating random numbers. Instead, they must simulate randomness, which is most often done using pseudo-random number generators.

A **pseudo-random number generator (PRNG)** is a program that takes a starting number (called a **seed**), and performs mathematical operations on it to transform it into some other number that appears to be unrelated to the seed. It then takes that generated number and performs the same mathematical operation on it to transform it into a new number that appears unrelated to the number it was generated from. By continually applying the algorithm to the last generated number, it can generate a series of new numbers that will appear to be random if the algorithm is complex enough.

Rule: You should only seed your random number generators once. Seeding them more than once will cause the results to be less random or non-random.

It's actually fairly easy to write a PRNG. Here's a short program that generates 100 pseudo-random numbers:

```

1  #include <iostream>
2
3  unsigned int PRNG()
4  {
5      // our initial starting seed is 5323
6      static unsigned int seed{ 5323 };
7
8      // Take the current seed and generate a new value from it
9      // Due to our use of large constants and overflow, it would be
10     // hard for someone to casually predict what the next number is
11     // going to be from the previous one.
12     seed = 8253729 * seed + 2396403;
13
14     // Take the seed and return a value between 0 and 32767
15     return seed % 32768;
16 }
17
18 int main()
19 {
20     // Print 100 random numbers
21     for (int count{ 1 }; count <= 100; ++count)
22     {
23         std::cout << PRNG() << '\t';

```

```

24
25     // If we've printed 5 numbers, start a new row
26     if (count % 5 == 0)
27         std::cout << '\n';
28     }
29
30     return 0;
31 }

```

The result of this program is:

```

23070  27857  22756  10839  27946
11613  30448  21987  22070  1001
27388  5999   5442   28789  13576
28411  10830  29441  21780  23687
5466   2957   19232  24595  22118
14873  5932   31135  28018  32421
14648  10539  23166  22833  12612
28343  7562   18877  32592  19011
13974  20553  9052   15311  9634
27861  7528   17243  27310  8033
28020  24807  1466   26605  4992
5235   30406  18041  3980   24063
15826  15109  24984  15755  23262
17809  2468   13079  19946  26141
1968   16035  5878   7337   23484
24623  13826  26933  1480   6075
11022  19393  1492   25927  30234
17485  23520  18643  5926   21209
2028   16991  3634   30565  2552
20971  23358  12785  25092  30583

```

Each number appears to be pretty random with respect to the previous one. As it turns out, our algorithm actually isn't very good, for reasons we will discuss later. But it does effectively illustrate the principle of PRNG number generation.

Generating random numbers in C++

C (and by extension C++) comes with a built-in pseudo-random number generator. It is implemented as two separate functions that live in the `cstdlib` header:

`std::srand()` sets the initial seed value to a value that is passed in by the caller. `srand()` should only be called once at the beginning of your program. This is usually done at the top of `main()`.

`std::rand()` generates the next random number in the sequence. That number will be a pseudo-random integer between 0 and `RAND_MAX`, a constant in `cstdlib` that is typically set to 32767.

Here's a sample program using these functions:

```

1  #include <iostream>
2  #include <cstdlib> // for std::rand() and std::srand()
3
4  int main()
5  {
6      std::srand(5323); // set initial seed value to 5323
7
8      // Print 100 random numbers
9      for (int count{ 1 }; count <= 100; ++count)

```

```

10     {
11         std::cout << std::rand() << '\t';
12
13         // If we've printed 5 numbers, start a new row
14         if (count % 5 == 0)
15             std::cout << '\n';
16     }
17
18     return 0;
19 }

```

Here's the output of this program:

```

17421    8558    19487    1344    26934
7796     28102  15201    17869    6911
4981     417    12650    28759    20778
31890    23714  29127    15819    29971
1069     25403  24427    9087     24392
15886    11466  15140    19801    14365
18458    18935  1746     16672    22281
16517    21847  27194    7163     13869
5923     27598  13463    15757    4520
15765    8582   23866    22389    29933
31607    180    17757    23924    31079
30105    23254  32726    11295    18712
29087    2787   4862     6569     6310
21221    28152  12539    5672     23344
28895    31278  21786    7674     15329
10307    16840  1645     15699    8401
22972    20731  24749    32505    29409
17906    11989  17051    32232    592
17312    32714  18411    17112    15510
8830     32592  25957    1269     6793

```

PRNG sequences and seeding

If you run the `std::rand()` sample program above multiple times, you will note that it prints the same result every time! This means that while each number in the sequence is seemingly random with regards to the previous ones, the entire sequence is not random at all! And that means our program ends up totally predictable (the same inputs lead to the same outputs every time). There are cases where this can be useful or even desired (e.g. you want a scientific simulation to be repeatable, or you're trying to debug why your random dungeon generator crashes).

But often, this is not what is desired. If you're writing a game of hi-lo (where the user has 10 tries to guess a number, and the computer tells them whether their guess is too high or too low), you don't want the program picking the same numbers each time. So let's take a deeper look at why this is happening, and how we can fix it.

Remember that each number in a PRNG sequence is generated from the previous number, in a deterministic way. Thus, given any starting seed number, PRNGs will always generate the same sequence of numbers from that seed as a result! We are getting the same sequence because our starting seed number is always 5323.

In order to make our entire sequence randomized, we need some way to pick a seed that's not a fixed number. The first answer that probably comes to mind is that we need a random number! That's a good thought, but if we need a random number to generate random numbers, then we're in a catch-22. It turns out, we really don't need our seed to be a random number -- we just need to pick something that changes each time the program is run. Then we can use our PRNG to generate a unique sequence of pseudo-random numbers from that seed.

The commonly accepted method for doing this is to enlist the system clock. Each time the user runs the program, the time will be different. If we use this time value as our seed, then our program will generate a different sequence of numbers each time it is run!

C comes with a function called `time()` that returns the number of seconds since midnight on Jan 1, 1970. To use it, we merely need to include the `ctime` header, and then initialize `srand()` with a call to `std::time(nullptr)` (or `std::time(0)` if your compiler is pre-C++11). We haven't covered `nullptr` yet, but it's essentially the equivalent of 0 in this context.

Here's the same program as above, using a call to `time()` as the seed:

```

1  #include <iostream>
2  #include <cstdlib> // for std::rand() and std::srand()
3  #include <ctime> // for std::time()
4
5  int main()
6  {
7      std::srand(static_cast<unsigned int>(std::time(nullptr))); // set initial seed value to sy
8
9      for (int count{ 1 }; count <= 100; ++count)
10     {
11         std::cout << std::rand() << '\t';
12
13         // If we've printed 5 numbers, start a new row
14         if (count % 5 == 0)
15             std::cout << '\n';
16     }
17
18     return 0;
19 }
```

Now our program will generate a different sequence of random numbers every time! Run it a couple of times and see for yourself.

Generating random numbers between two arbitrary values

Generally, we do not want random numbers between 0 and `RAND_MAX` -- we want numbers between two other values, which we'll call `min` and `max`. For example, if we're trying to simulate the user rolling a die, we want random numbers between 1 and 6 (pedantic grammar note: yes, die is the singular of dice).

Here's a short function that converts the result of `rand()` into the range we want:

```

1  // Generate a random number between min and max (inclusive)
2  // Assumes std::srand() has already been called
3  // Assumes max - min <= RAND_MAX
4  int getRandomNumber(int min, int max)
5  {
6      static constexpr double fraction { 1.0 / (RAND_MAX + 1.0) }; // static used for efficiency
7      // evenly distribute the random number across our range
8      return min + static_cast<int>((max - min + 1) * (std::rand() * fraction));
9  }
```

To simulate the roll of a die, we'd call `getRandomNumber(1, 6)`. To pick a randomized digit, we'd call `getRandomNumber(0, 9)`.

Optional reading: How does the previous function work?

The `getRandomNumber()` function may seem a little complicated, but it's not too bad.

Let's revisit our goal. The function `rand()` returns a number between 0 and `RAND_MAX` (inclusive). We want to somehow transform the result of `rand()` into a number between `min` and `max` (inclusive). This means that when we

do our transformation, 0 should become min, and RAND_MAX should become max, with a uniform distribution of numbers in between.

We do that in five parts:

1. We multiply our result from `std::rand()` by fraction. This converts the result of `rand()` to a floating point number between 0 (inclusive), and 1 (exclusive).

If `rand()` returns a 0, then $0 * \text{fraction}$ is still 0. If `rand()` returns `RAND_MAX`, then $\text{RAND_MAX} * \text{fraction}$ is $\text{RAND_MAX} / (\text{RAND_MAX} + 1)$, which is slightly less than 1. Any other number returned by `rand()` will be evenly distributed between these two points.
2. Next, we need to know how many numbers we can possibly return. In other words, how many numbers are between min (inclusive) and max (inclusive)?

This is simply $(\text{max} - \text{min} + 1)$. For example, if `max = 8` and `min = 5`, $(\text{max} - \text{min} + 1) = (8 - 5 + 1) = 4$. There are 4 numbers between 5 and 8 (that is, 5, 6, 7, and 8).
3. We multiply the prior two results together. If we had a floating point number between 0 (inclusive) and 1 (exclusive), and then we multiply by $(\text{max} - \text{min} + 1)$, we now have a floating point number between 0 (inclusive) and $(\text{max} - \text{min} + 1)$ (exclusive).
4. We cast the previous result to an integer. This removes any fractional component, leaving us with an integer result between 0 (inclusive) and $(\text{max} - \text{min})$ (inclusive).
5. Finally, we add min, which shifts our result to an integer between min (inclusive) and max (inclusive).

Optional reading: Why don't we use the modulus operator (%) in the previous function?

One of the most common questions readers have submitted is why we use division in the above function instead of modulus (%). The short answer is that the modulus method tends to be biased in favor of low numbers.

Let's consider what would happen if the above function looked like this instead:

```
1 | return min + (std::rand() % (max-min+1));
```

Seems similar, right? Let's explore where this goes wrong. To simplify the example, let's say that `rand()` always returns a random number between 0 and 9 (inclusive). For our sample case, we'll pick `min = 0`, and `max = 6`. Thus, $\text{max} - \text{min} + 1$ is 7.

Now let's calculate all possible outcomes:

```
0 + (0 % 7) = 0
0 + (1 % 7) = 1
0 + (2 % 7) = 2
0 + (3 % 7) = 3
0 + (4 % 7) = 4
0 + (5 % 7) = 5
0 + (6 % 7) = 6

0 + (7 % 7) = 0
0 + (8 % 7) = 1
0 + (9 % 7) = 2
```

Look at the distribution of results. The results 0 through 2 come up twice, whereas 3 through 6 come up only once. This method has a clear bias towards low results. By extension, most cases involving this algorithm will behave similarly.

Now let's take a look at the result of the `getRandomNumber()` function above, using the same parameters as above (`rand()` returns a number between 0 and 9 (inclusive), `min = 0` and `max = 6`). In this case, $\text{fraction} = 1 / (9 + 1) = 0.1$.

max - min + 1 is still 7.

Calculating all possible outcomes:

```
0 + static_cast(7 * (0 * 0.1)) = 0 + static_cast(0) = 0
0 + static_cast(7 * (1 * 0.1)) = 0 + static_cast(0.7) = 0
0 + static_cast(7 * (2 * 0.1)) = 0 + static_cast(1.4) = 1
0 + static_cast(7 * (3 * 0.1)) = 0 + static_cast(2.1) = 2
0 + static_cast(7 * (4 * 0.1)) = 0 + static_cast(2.8) = 2
0 + static_cast(7 * (5 * 0.1)) = 0 + static_cast(3.5) = 3
0 + static_cast(7 * (6 * 0.1)) = 0 + static_cast(4.2) = 4
0 + static_cast(7 * (7 * 0.1)) = 0 + static_cast(4.9) = 4
0 + static_cast(7 * (8 * 0.1)) = 0 + static_cast(5.6) = 5
0 + static_cast(7 * (9 * 0.1)) = 0 + static_cast(6.3) = 6
```

The bias here is still slightly towards lower numbers (0, 2, and 4 appear twice, whereas 1, 3, 5, and 6 appear once), but it's much more uniformly distributed.

Even though `getRandomNumber()` is a little more complicated to understand than the modulus alternative, we advocate for the division method because it produces a less biased result.

What is a good PRNG?

As I mentioned above, the PRNG we wrote isn't a very good one. This section will discuss the reasons why. It is optional reading because it's not strictly related to C or C++, but if you like programming you will probably find it interesting anyway.

In order to be a good PRNG, the PRNG needs to exhibit a number of properties:

First, the PRNG should generate each number with approximately the same probability. This is called distribution uniformity. If some numbers are generated more often than others, the result of the program that uses the PRNG will be biased!

For example, let's say you're trying to write a random item generator for a game. You'll pick a random number between 1 and 10, and if the result is a 10, the monster will drop a powerful item instead of a common one. You would expect a 1 in 10 chance of this happening. But if the underlying PRNG is not uniform, and generates a lot more 10s than it should, your players will end up getting more rare items than you'd intended, possibly trivializing the difficulty of your game.

Generating PRNGs that produce uniform results is difficult, and it's one of the main reasons the PRNG we wrote at the top of this lesson isn't a very good PRNG.

Second, the method by which the next number in the sequence is generated shouldn't be obvious or predictable. For example, consider the following PRNG algorithm: `num = num + 1`. This PRNG is perfectly uniform, but it's not very useful as a sequence of random numbers!

Third, the PRNG should have a good dimensional distribution of numbers. This means it should return low numbers, middle numbers, and high numbers seemingly at random. A PRNG that returned all low numbers, then all high numbers may be uniform and non-predictable, but it's still going to lead to biased results, particularly if the number of random numbers you actually use is small.

Fourth, all PRNGs are periodic, which means that at some point the sequence of numbers generated will eventually begin to repeat itself. As mentioned before, PRNGs are deterministic, and given an input number, a PRNG will produce the same output number every time. Consider what happens when a PRNG generates a number it has previously generated. From that point forward, it will begin to duplicate the sequence between the first occurrence of that number and the next occurrence of that number over and over. The length of this sequence is known as the **period**.

For example, here are the first 100 numbers generated from a PRNG with poor periodicity:

| | | | | |
|-----|-----|-----|-----|-----|
| 112 | 9 | 130 | 97 | 64 |
| 31 | 152 | 119 | 86 | 53 |
| 20 | 141 | 108 | 75 | 42 |
| 9 | 130 | 97 | 64 | 31 |
| 152 | 119 | 86 | 53 | 20 |
| 141 | 108 | 75 | 42 | 9 |
| 130 | 97 | 64 | 31 | 152 |
| 119 | 86 | 53 | 20 | 141 |
| 108 | 75 | 42 | 9 | 130 |
| 97 | 64 | 31 | 152 | 119 |
| 86 | 53 | 20 | 141 | 108 |
| 75 | 42 | 9 | 130 | 97 |
| 64 | 31 | 152 | 119 | 86 |
| 53 | 20 | 141 | 108 | 75 |
| 42 | 9 | 130 | 97 | 64 |
| 31 | 152 | 119 | 86 | 53 |
| 20 | 141 | 108 | 75 | 42 |
| 9 | 130 | 97 | 64 | 31 |
| 152 | 119 | 86 | 53 | 20 |
| 141 | 108 | 75 | 42 | 9 |

You will note that it generated 9 as the second number, and 9 again as the 16th number. The PRNG gets stuck generating the sequence in-between these two 9's repeatedly: 9-130-97-64-31-152-119-86-53-20-141-108-75-42-(repeat).

A good PRNG should have a long period for *all* seed numbers. Designing an algorithm that meets this property can be extremely difficult -- most PRNGs will have long periods for some seeds and short periods for others. If the user happens to pick a seed that has a short period, then the PRNG won't be doing a good job.

Despite the difficulty in designing algorithms that meet all of these criteria, a lot of research has been done in this area because of its importance to scientific computing.

std::rand() is a mediocre PRNG

The algorithm used to implement `std::rand()` can vary from compiler to compiler, leading to results that may not be consistent across compilers. Most implementations of `rand()` use a method called a **Linear Congruential Generator (LCG)**. If you have a look at the first example in this lesson, you'll note that it's actually a LCG, though one with intentionally picked poor constants. LCGs tend to have shortcomings that make them not good choices for most kinds of problems.

One of the main shortcomings of `rand()` is that `RAND_MAX` is usually set to 32767 (essentially 15-bits). This means if you want to generate numbers over a larger range (e.g. 32-bit integers), `rand()` is not suitable. Also, `rand()` isn't good if you want to generate random floating point numbers (e.g. between 0.0 and 1.0), which is often useful when doing statistical modelling. Finally, `rand()` tends to have a relatively short period compared to other algorithms.

That said, `rand()` is perfectly suitable for learning how to program, and for programs in which a high-quality PRNG is not a necessity.

For applications where a high-quality PRNG is useful, I would recommend **Mersenne Twister** (or one of its variants), which produces great results and is relatively easy to use. Mersenne Twister was adopted into C++11, and we'll show how to use it later in this lesson.

Debugging programs that use random numbers

Programs that use random numbers can be difficult to debug because the program may exhibit different behaviors each time it is run. Sometimes it may work, and sometimes it may not. When debugging, it's helpful to ensure your program executes the same (incorrect) way each time. That way, you can run the program as many times as needed to isolate where the error is.

For this reason, when debugging, it's a useful technique to set the random seed (via `std::srand`) to a specific value (e.g. 0) that causes the erroneous behavior to occur. This will ensure your program generates the same results each time, making debugging easier. Once you've found the error, you can seed using the system clock again to start generating randomized results again.

Random numbers in C++11

C++11 added a ton of random number generation functionality to the C++ standard library, including the Mersenne Twister algorithm, as well as generators for different kinds of random distributions (uniform, normal, Poisson, etc...). This is accessed via the `<random>` header.

Here's a short example showing how to generate random numbers in C++11 using Mersenne Twister (h/t to user Fernando):

```

1  #include <iostream>
2  #include <random> // for std::mt19937
3  #include <ctime> // for std::time
4
5  int main()
6  {
7      // Initialize our mersenne twister with a random seed based on the clock
8      std::mt19937 mersenne{ static_cast<std::mt19937::result_type>(std::time(nullptr)) };
9
10     // Create a reusable random number generator that generates uniform numbers between 1 and
11     std::uniform_int_distribution die{ 1, 6 };
12     // If your compiler doesn't support C++17, use this instead
13     // std::uniform_int_distribution<> die{ 1, 6 };
14
15     // Print a bunch of random numbers
16     for (int count{ 1 }; count <= 48; ++count)
17     {
18         std::cout << die(mersenne) << '\t'; // generate a roll of the die here
19
20         // If we've printed 6 numbers, start a new row
21         if (count % 6 == 0)
22             std::cout << '\n';
23     }
24
25     return 0;
26 }
```

Author's note

Before C++17, you need to use add empty brackets to create die after the type `std::uniform_int_distribution<> die{ 1, 6 }`

You'll note that Mersenne Twister generates random 32-bit unsigned integers (not 15-bit integers like `std::rand()`), giving a lot more range. There's also a version (`std::mt19937_64`) for generating 64-bit unsigned integers.

C++11 style random numbers across multiple functions

The above example create a random generator for use within a single function. What happens if we want to use a random number generator in multiple functions?

Although you can create a static local `std::mt19937` variable in each function that needs it (static so that it only gets seeded once), it's a little overkill to have every function that needs a random number generator seed and maintain its own local generator. A better option in most cases is to create a global random number generator (inside a namespace!). Remember how we told you to avoid non-const global variables? This is an exception (also note: `std::rand()` and `std::srand()` access a global object, so there's precedent for this).

```

1  #include <iostream>
2  #include <random> // for std::mt19937
3  #include <ctime> // for std::time
4
5  namespace MyRandom
6  {
7      // Initialize our mersenne twister with a random seed based on the clock (once at system s
8      std::mt19937 mersenne{ static_cast<std::mt19937::result_type>(std::time(nullptr)) };
9  }
10
11 int getRandomNumber(int min, int max)
12 {
13     std::uniform_int_distribution die{ min, max }; // we can create a distribution in any func
14     return die(MyRandom::mersenne); // and then generate a random number from our global gener
15 }
16
17 int main()
18 {
19     std::cout << getRandomNumber(1, 6) << '\n';
20     std::cout << getRandomNumber(1, 10) << '\n';
21     std::cout << getRandomNumber(1, 20) << '\n';
22
23     return 0;
24 }
```

Using a random library

A perhaps better solution is to use a 3rd party library that handles all of this stuff for you, such as the header-only **Effolkronium's random library**. You simply add the header to your project, `#include` it, and then you can start generating random numbers via `Random::get(min, max)`.

Here's the above program using Effolkronium's library:

```

1  #include <iostream>
2  #include "random.hpp"
3
4  // get base random alias which is auto seeded and has static API and internal state
5  using Random = effolkronium::random_static;
6
7  int main()
8  {
9      std::cout << Random::get(1, 6) << '\n';
10     std::cout << Random::get(1, 10) << '\n';
11     std::cout << Random::get(1, 20) << '\n';
12
13     return 0;
14 }
```

Help! My random number generator is generating the same sequence of random numbers!

If your random number generator is generating the same sequence of random numbers every time your program is run, you probably didn't seed it properly. Make sure you're seeding it with a value that changes each time the program is run (like `std::time(nullptr)`).

Help! My random number generator always generates the same first number!

The implementation of `rand()` in Visual Studio and a few other compilers has a flaw -- the first random number generated doesn't change much for similar seed values. This means that when using `std::time(nullptr)` to seed your random number generator, the first result from `rand()` won't change much in successive runs. However, the results of successive calls to `rand()` aren't impacted, and will be sufficiently randomized.

The solution here, and a good rule of thumb in general, is to discard the first random number generated from the random number generator.

Help! My random number generator isn't generating random numbers at all!

If your random number generator is generating the same number every time you ask it for a random number, then you are probably either reseeding the random number generator before generating a random number, or you're creating a new random generator for each random number.

Here's are two functions that exhibit the issue:

```

1  int rollDie()
2  {
3      std::srand(std::time(nullptr));
4      return getRandomNumber(1, 6); // using definition of getRandomNumber above
5  }
6
7  int getRandomNumber()
8  {
9      std::mt19937 mersenne{ static_cast<std::mt19937::result_type>(std::time(nullptr)) };
10     std::uniform_int_distribution rand{ 1, 52 };
11     return rand(mersenne);
12 }
```

In both cases, the random number generator is being seeded each time before a random number is generated. This will cause a similar number to be generated each time.

In the top case, `std::srand()` is reseeding the built-in random number generator before `rand()` is called (by `getRandomNumber()`).

In the bottom case, we're creating a new Mersenne Twister, seeding it, generating a single random number, and then destroying it.

For best results, you should only seed a random number generator once (generally at program initialization for `std::srand()`, or the point of creation for other random number generators), and then use that same random number generator for each successive random number generated.



[5.10 -- std::cin, extraction, and dealing with invalid text input](#)



[Index](#)



[5.8 -- Break and continue](#)

 C++ TUTORIAL, GAME PROGRAMMING
  C++, PRNG, PROGRAMMING, RAND, RANDOM, RNG, SRAND, TUTORIAL
 |  PRINT THIS POST

443 comments to 5.9 — Random number generation

[« Older Comments](#)
1
...
4
5
6



Charan

December 26, 2019 at 9:39 pm · Reply

Hey, How do I get effolkronium's library? My compiler through a 'fatal error' saying "random.hpp: no such file or directory".



kavin

January 16, 2020 at 10:39 am · Reply

Same error for me too. I think u have to include it via github or something like that. I guess it's beyond the scope of this lesson.



chai

December 23, 2019 at 11:20 am · Reply

```

1  int main()
2  {
3      // Initialize our mersenne twister with a random seed based on the clock
4      std::mt19937 mersenne{ static_cast<std::mt19937::result_type>(std::time(nullptr)) };
5
6      // Create a reusable random number generator that generates uniform numbers between 1 a
7      std::uniform_int_distribution die{ 1, 6 }; // Argument for class template is missing
8
9      // Print a bunch of random numbers
10     for (int count{ 1 }; count <= 48; ++count)
11     {
12         std::cout << die(mersenne) << '\t'; // Identifier die is undefined
13
14         // If we've printed 6 numbers, start a new row
15         if (count % 6 == 0)

```

```

16         std::cout << '\n';
17     }
18
19     return 0;
20 }
```

There are two red squiggly lines that showed up in my Visual Studio after cutting and pasting the example code. One said that:- Argument for class template "std::uniform_int_distribution" is missing.

The other is "die" is undefined.

Any idea how to correct this?



nascar driver

December 24, 2019 at 2:52 am · Reply

The error is correct if you didn't enable C++17 or higher. Enable C++17 or higher in your project settings. If you can't, change line 7 to

```

1 //
2 std::uniform_int_distribution<> die{ 1, 6 };
```



chai

December 24, 2019 at 8:06 am · Reply

awesome! Thanks.



Ged

December 13, 2019 at 9:23 am · Reply

Just finished everything till OOP. So came back to RNG to ask a few questions.

1. When we use the library `<cstdlib>`, it declares our `RAND_MAX` as a MACRO to be typically a value of 32767, so that is why we don't need to declare it ourselves?
2. Why do we need to use `static_cast<int>` if we are returning an int, won't it chop off everything when it returns the value?

```

1 int getRandomNumber(int min, int max)
2 {
3     static constexpr double fraction{ 1.0 / (RAND_MAX + 1.0) };
4     return min + static_cast<int>((max - min + 1) * (std::rand() * fraction));
5 }
```

3. Understanding the last RNG (Mersenne Twister) is a bit hard and as you mentioned in order to fully understand it you need to know classes and templates. So not understanding it fully is ok for now?
4. Is it even worth trying to write your own RNG or it's easier to download 3rd party libraries that requires you to write a single line for a good RNG? If a 3rd party client will always be a better option then what if you are writing a complex program, will you need a lot of 3rd party libraries in order for your code to be much shorter?
5. Is there a chapter in the future that will show us how to import the libraries into your program?



nascar driver

December 13, 2019 at 9:44 am · Reply

1. Yes.

2. Yes, the code does the same without the cast. There might be the unlikely event that double is implemented in a way such that it can't hold all values an int can hold. In that case, the two versions might behave differently.
3. Yes.
4. Use the standard functions unless you need something specialized.
5. Appendix A. The process differs between compilers and IDEs. If yours isn't in the appendix, google will find instructions.



Elis

November 13, 2019 at 4:34 am · Reply

As part of a "Hi-lo game" I used the following randomization process. Just using the randomize function would always yield the same result in lower ranges. E.g. loNumber == 3 and hiNumber == 9 would always yield 7 e.t.c.
 With this I managed to get pretty good results.
 Since this runs the randomize function many times in a row is it a necessity to place the seed elsewhere than in the function itself (as below)?

```

1  int randomize(int hiNumber, int loNumber)
2  {
3      static const double fraction{ 1.0 / (RAND_MAX + 1.0) };
4      int random = (loNumber + static_cast<int>((hiNumber - loNumber + 1) * (std::rand() * fr
5      return random;
6  }
7
8  int generateNumber(int hiNumber, int loNumber)
9  {
10     std::srand(static_cast<int>(std::time(nullptr)));           //seed
11     int random{ randomize(hiNumber, loNumber) };               //random number for tim
12     int returnValue{};
13     for (int count{ 1 }; count <= ((random % 50) + 5); ++count) /*get a random va
14         random nr of times(between 5 and 50)*/
15         returnValue = randomize(hiNumber, loNumber);
16
17     return returnValue;
18 }
```

As always thanks for the great tutorials!



nascar driver

November 13, 2019 at 4:38 am · Reply

Yes. Never seed more than once. Seeding again with the same seed will reset the sequence, so you'll get the same "random" number again.



Ged

November 6, 2019 at 11:08 am · Reply

I don't understand these lines fully. Would appreciate if you could explain a bit more.

```

1  // Initialize our mersenne twister with a random seed based on the clock
2  std::mt19937 mersenne(static_cast<std::mt19937::result_type>(std::time(nullptr))); // Wh
3
4  // Create a reusable random number generator that generates uniform numbers between 1 an
5  std::uniform_int_distribution<> die(1, 6); // how does uniform_int_distribution even wor
```



nascar driver

November 7, 2019 at 2:14 am · Reply

You won't be able to understand them fully until you get to classes and templates. `std::my19937` wants a `std::mt19937::result_type`, but `std::time` returns an unspecified type, so we need to cast it.

> how does `uniform_int_distribution` even work

There's no general answer, as the standard only says `_what_`std::uniform_int_distribution` does`, but not `_how_`. If you want to understand how it works, you'll have to read up on discrete probability, finish learning C++, and then try to read a standard implementation.

> why do we have `<>`

They're not necessary since C++17, you can write `std::uniform_int_distribution die{ 1, 6 };``. Inside the `<>`, you can specify the integer type you want to have as a result, eg. `std::uniform_int_distribution die{ 1, 6 };``

> how does `die(1,6)`

You'll need to know about classes and constructors. Basically, `die`` stores the numbers you give it (min and max) and uses them later when it generates a number.



Tushar

November 3, 2019 at 9:02 am · Reply

```

1  #include<iostream>
2  unsigned int PRNG(){
3
4      static unsigned int seed=9843;
5      seed= seed*8343454 + 3555643;
6
7      return seed%34567;
8  }
9  int main(){
10     for(int i=1;i<100;i++){
11         std::cout<<PRNG()<<"\t";
12         if(i%5==0){
13             std::cout<<"\n";
14         }
15     }
16
17 }
```

Output:

```

990  3213  1565  28503  21640
5362  29035  4686  32618  27599
7500  3776  13427  15361  8288
30298  17693  11683  32746  9669
10489  9238  19735  17318  20090
18670  12737  24603  871  13768
22541  22541  22541  22541  22541
22541  22541  22541  22541  22541
22541  22541  22541  22541  22541
22541  22541  22541  22541  22541
22541  22541  22541  22541  22541
```

```

22541 22541 22541 22541 22541
22541 22541 22541 22541 22541
22541 22541 22541 22541 22541
22541 22541 22541 22541 22541
22541 22541 22541 22541 22541
22541 22541 22541 22541 22541
22541 22541 22541 22541 22541
22541 22541 22541 22541 22541
22541 22541 22541 22541

```

Why are the numbers being repeated?



Luiz

November 3, 2019 at 12:17 pm · Reply.

It's going to take someone more experienced than me to explain but I managed to fix it by doing the following:

snip

```

1  unsigned int PRNG() {
2
3      static unsigned int seed = 9843;
4      seed = seed * 8343454 + 3555643;
5      seed %= 34567; // Seed wasn't being assigned to it's module
6      return seed;
7  }

```

I assigned seed to it's module by 34567 everytime

It's probably got to do with undefined behaviour via overflow



Tushar

November 3, 2019 at 7:49 pm · Reply.

Okay so this has reduced repetition of a single number though its producing repeated pattern

Output:

```

3381 2929 3321 745 1801
2633 1993 3401 2505 1737
1225 2889 713 1097 73
3401 2505 1737 1225 2889
713 1097 73 3401 2505
1737 1225 2889 713 1097
73 3401 2505 1737 1225
2889 713 1097 73 3401
2505 1737 1225 2889 713
1097 73 3401 2505 1737
1225 2889 713 1097 73
3401 2505 1737 1225 2889
713 1097 73 3401 2505
1737 1225 2889 713 1097
73 3401 2505 1737 1225
2889 713 1097 73 3401
2505 1737 1225 2889 713
1097 73 3401 2505 1737

```

1225 2889 713 1097 73

3401 2505 1737 1225 2889

You can see the pattern starting with 713 1097 73... 2889.



Luiz

[November 4, 2019 at 3:12 am · Reply](#)

There's no error within the code, it is all caused by the math and overflowing of the unsigned int. The operations your number go through in your pseudo random generator have created a pattern. If you want to try to create one that's less predictable you could try adding more operations of different sorts or such. I reckon this is also true for the reason as to why 22541 was being repeated on the first version of your code.



Luiz

[November 4, 2019 at 3:16 am · Reply](#)

Also as it is said in this lesson: The PRNG is perfectly uniform, so you're better off using a different method to generate random numbers



Tushar

[November 5, 2019 at 1:21 am · Reply](#)

```

1  #include<iostream>
2  #include<cstdlib>
3  #include<time.h>
4
5  unsigned int PRNG(int start,int end,int exceed){
6      unsigned int seed=6784;
7      for(int i=start;i<end;i++){
8          seed+=seed /end *(time(0)/1000) + exceed;
9          std::cout<<seed /100000<<"\t";
10         if(i%5==0){
11             std::cout<<"\n\n";
12         }
13     }
14     return seed;
15 }
16 int main(){
17
18     PRNG(1,100,10);
19 }
```

Doesn't generate repeated number nor a pattern!
Is it correct?



nascardriver

[November 5, 2019 at 2:32 am · Reply](#)

`end` and `std::time(nullptr)` are constant, you're basically doing

```
1 | seed += ((seed * k) + exceed);
```


just more expensive. The modulus that you removed is now performed by the wrap-around of `seed`.

Take your original code and change the constant factor to an odd number, that should get rid of the repetition.

```

1  #include <iostream>
2
3  unsigned int PRNG()
4  {
5      static unsigned int seed{ 9843 };
6      seed = seed * 8343453 + 3555643;
7      //          ^^^^^^^ odd
8
9      return seed % 34567;
10 }
11
12 int main()
13 {
14     for (int i{ 1 }; i <= 1000; ++i)
15     {
16         std::cout << PRNG() << '\t';
17         if ((i % 5) == 0)
18         {
19             std::cout << "\n";
20         }
21     }
22
23     return 0;
24 }
```

I don't know why this works, Mike observed something similar here <https://www.learncpp.com/cpp-tutorial/59-random-number-generation/comment-page-6/#comment-427692>



Mike

October 16, 2019 at 3:18 pm · Reply.

"Consider what happens when a PRNG generates a number it has previously generated. From that point forward, it will begin to duplicate the sequence between the first occurrence of that number and the next occurrence of that number over and over. The length of this sequence is known as the period."

Is this behavior only with large ranges? Because when I do it with a range of 1-6 (roll of die) or other similar low ranges, the numbers often get repeated many times, yet I haven't noticed a pattern from any of them. Based on what Alex stated above, shouldn't a pattern start right after the first number is repeated?

Other than the much larger range, how are these different than the example that Alex provided?



nascar driver

October 17, 2019 at 1:04 am · Reply.

When you generate a number from 1 to 6, you're using a way larger number and scale it down. When you see two 1s, they're probably not both based on the same number.

Consider a simple `std::rand` generator

```
1 | int die{ (std::rand() % 6) + 1 };
```

`die` is 1 for 1, 7, 13, 19, etc.



Mike

[October 17, 2019 at 9:25 am · Reply](#)

Thanks for the feedback!

So if I understood you correctly, scaling down from a larger number, like maybe the RAND_MAX of 32767, to a significantly smaller range of numbers, like the roll of a die, will likely contribute to a longer or better period?

**nascar driver**[October 18, 2019 at 1:27 am · Reply](#)

Scaling down the number doesn't affect the period. The period is caused by the implementation of the PRNG, which you're not changing.



hellmet

[October 16, 2019 at 5:24 am · Reply](#)

I have a few questions with the examples here.

```

1  ...
2      Explain what's happening here please...
3      std::mt19937 mersenne(static_cast<std::mt19937::result_type>(std::time(nullptr)));
4
5      std::uniform_int_distribution<> die(1, 6); <- What is die here? Looks like variabl
6
7      // Print a bunch of random numbers
8      for (int count = 1; count <= 48; ++count)
9      {
10         std::cout << die(mersenne) << "\t"; <- die used like a function here?
11     }
12     ...

```

What is 'mersenne' here? I think it's the uniform initialization of

What is the type of 'die' here? It looks like a variable but later used as a function?

**nascar driver**[October 16, 2019 at 5:34 am · Reply](#)

`mersenne` is an `std::mt19937`, a mersenne twister. It's a random number generator, seeded with the current time.

`die`'s type is `std::uniform_int_distribution`, you don't need the `<>`. It's a variable, but its `operator()` has been overloaded so it can be used as if it was a function (Covered later).

`die` uses `mersenne` to generate a random number. Then it turns the random number into a number in the range of 1 to 6 with all numbers being equally likely to occur.



hellmet

[October 16, 2019 at 7:05 am · Reply](#)

Ahh I see, variable with () is an overloaded operator, cool I'll get back to this page once that is covered; and the Mersenne is a variable that's being uniform initialized with time that's being static cast to a std::mt199...::result type. Got it, Thanks a lot!

nascar driver



October 16, 2019 at 7:30 am · Reply.

> the Mersenne is a variable that's being uniform initialized
It's not uniform initialized. Uniform initialization (Brace initialization) uses curly braces

```
1 | int i{ 123 }; // Brace initialization (Type safe)
2 | int i(123); // Direct initialization (Not so type safe).
```

Brace initialization is a form of direct initialization, but with higher type safety and it can be used to initialize lists.



hellmet

October 16, 2019 at 10:21 am · Reply.

My bad! Got the names mixed up. So kind of you point that out too! Thank you!



Jon


October 7, 2019 at 2:32 am · Reply.

I hope you don't mind a couple of comments related to pedagogy...

I wonder if it's worth mentioning, at least in passing, cryptographically strong random and why you shouldn't use that for anything but crypto.

My take is that this lesson, while quite good, is a big jump from the last one where we just learned about the for loop. We still haven't been taught about angle brackets and such. Personally, I'm not bothered by the pacing here but I've been programming for decades and I'm just brushing up on modern C++.

Thank you so much for the lessons and please don't take this as criticism, it's just food for thought.



Alex

October 8, 2019 at 4:00 pm · Reply.

I never mind feedback, criticism or otherwise. It's part of how we learn and grow.

> I wonder if it's worth mentioning, at least in passing, cryptographically strong random and why you shouldn't use that for anything but crypto.

Curious why this is relevant to the lesson.

> My take is that this lesson, while quite good, is a big jump from the last one where we just learned about the for loop

What parts do you find to be big jumps? The heavy algorithmic focus, or something more language-syntax specific?

> We still haven't been taught about angle brackets and such

Somewhere prior to this point I mention that angle brackets are how we parameterize types. I'll see if I can find and reinforce that, because it's an important point prior to being able to describe how they actually function (covered in the lessons on templates).


Mike

October 1, 2019 at 11:27 am · Reply.



When debugging, how do you watch or monitor a value that keeps changing, yet doesn't have a variable assigned to it, like this PRNG statement that I took from a For Loop? Obviously, the value is there and it keeps changing, but even hovering your mouse over it while debugging will not display any results.

```
std::cout << std::rand() << '\t';
```



Alex

October 2, 2019 at 4:41 pm · Reply.

One thing that works sometimes is placing a breakpoint on the entry point to wherever the intermediate value (e.g. the result of `std::rand()`) is being sent. In this case, that would be to `operator<<`, which is likely going to hit everything.

Probably the easiest thing to do is assign the value you're interested in to a variable so it persists beyond the life of the expression, and examine it immediately after.

```
1 | int temp;
2 | std::cout << (temp = std::rand()) << '\t';
```

Anybody else have any better tips here?



nascar driver

October 3, 2019 at 12:36 am · Reply.

I didn't reply because this depends on the debugger and IDE and I suppose Mike doesn't use gdb.

In gdb, you can "break" on the `std::cout ...` line, "step" into `std::rand`, then "finish" and it prints the return value.

I've seen an integrated debugger that shows return values, but I don't think it was for cpp. Though I'm sure one exist for cpp as well.

If the source code can be modified, the line can be split

```
1 | std::cout
2 |     << std::rand()
3 |     << '\t';
```

then break on the `\t` line and "print \$rax" if the architecture uses rax for return values and it's an integer.

Personally I do

```
1 | int temp{ std::rand() };
2 |
3 | // Debugging without a debugger.
4 | std::cout << temp << '\n';
5 |
6 | // Doesn't make a lot of sense with this example, imagine this
7 | // line was something else
8 | std::cout << temp << '\n';
```



Mike

October 1, 2019 at 10:46 am · Reply.

I notice that if I alter the following line in the first code snippet from Ch5.9, that after 13 results, the remaining 87 results are always 32767. At first, I was confused by this, but then I realized, it must be because during each iteration of `@PRNG()`, the `@seed` value is always increasing. So at some point, the

@seed value is greater than 32768, and so the difference is wrapped. But at some point the amount that is now wrapped is even larger than the 32768, which apparently it can't wrap around again, and thus keeps displaying 32768. So does this mean values can only wrap around once, or is something else at play here?

My primary question is though, why when using the seed values Alex provided in the code snippet, this behavior does not happen? And his numbers are much larger than the one's I used. I even changed the count to 1000, yet the 32768 never repeats. Is he using some special magic numbers to make this happen?

original code from Ch5.9, first snippet:

```
seed = 8253729 * seed + 2396403;
```

my alteration:

```
seed = 2 * seed + 1;
```



nascar driver

October 2, 2019 at 4:33 am · Reply

The multiplication by 2 is the problem. Once you generate `seed - 1`, you'll always get `seed - 1` as the result, because

```
1 // @a is the previous seed
2 // @n is what we modulate by (In the quiz 32768).
3 // we know
4 (a % n) = (n - 1)
5
6 // Now we update the seed (a = 2a + 1).
7 ((2a + 1) % n)
8 = ((2a % n) + (1 % n)) % n
9 = (((2a % n) + 1) % n)
10 = (((2 * (a % n)) % n) + 1) % n
11 = (((2 * (n - 1)) % n) + 1) % n
12 = ((n - 2) + 1) % n
13 = (n - 1) % n
14 = (n - 1) // we're back where we started
```

You'll observe the same behavior when you replace 2 with any even number (I didn't prove that, but seems to work out).

> using the seed values Alex provided in the code snippet, this behavior does not happen
He uses an odd number for the multiplication. Add or subtract one and it'll get stuck too.



Mike

October 2, 2019 at 7:14 am · Reply

Interesting, it makes since now. It's also good to know the additional wrapping was not the issue, though for a nascent learner like myself, it seemed logical.




SullenSecret

September 17, 2019 at 10:07 pm · Reply

I'm sorry, but I feel like many parts of this lesson need further explanations. I had more trouble with it than usual due to certain details simply not being explained. It's as if this topic is disliked by the author.

Alex

September 19, 2019 at 4:02 pm · Reply



Can you be more specific about which parts were hard for you to understand? Thanks!



Raiyan

September 7, 2019 at 2:49 pm · Reply

What's the difference between

```
1 | std::mt19937 mersenne { static_cast<std::mt19937::result_type>(std::time(nullptr)) }
```

and

```
1 | std::mt19937 mersenne { static_cast<unsigned int>(std::time(nullptr)) }
```

.



nascardriver

September 8, 2019 at 12:48 am · Reply

`std::mt19937::result_type` is an `std::uint_fast32_t`, a 32-bit wide integer.

`unsigned int` is at least 16 bits wide.

By casting to `unsigned int` you're discarding 16 bits, greatly decreasing the number of potential seeds. On top of that, `std::mt19937` doesn't want to have an `unsigned int`, so an implicit cast has to be performed to convert the `unsigned int` to `std::mt19937::result_type`.

Using `unsigned int` here is unfounded.



alfonso

August 18, 2019 at 1:38 am · Reply

```
1 | #include <iostream> // for standard I/O
2 | #include <ctime> // for std::time ()
3 | #include <random> // for mersenne twister
4 |
5 | namespace MyRandom {
6 |     std::mt19937_64 seed {
7 |         static_cast <std::mt19937_64::result_type> (std::time (nullptr))
8 |     };
9 | }
10 |
11 | double randomReal (double min, double max) {
12 |     std::mt19937_64 seed {
13 |         static_cast <std::mt19937_64::result_type> (std::time (nullptr))
14 |     };
15 |     std::uniform_real_distribution <> realNumber (min, max);
16 |
17 |     return realNumber (seed); // CASE 1
18 |     // return realNumber (MyRandom::seed); // CASE 2
19 | }
20 |
21 | int main () {
22 |     std::cout << randomReal (1.0, 3.0) << '\n';
23 |     std::cout << randomReal (1.0, 3.0) << '\n';
24 |     std::cout << randomReal (1.0, 3.0) << '\n';
25 |     return 0;
26 | }
```

I do not understand why. If I use CASE 1, the returned "random" number is always the same. But if I use CASE 2 (with global (namespaced) seed) the result is as expected, random.



nascardriver

August 18, 2019 at 1:47 am · Reply.

See the paragraph "Help! My random number generator isn't generating random numbers at all!" at the end of this lesson.



alfonso

August 18, 2019 at 11:53 pm · Reply.

"The random number generator is being seeded each time before a random number is generated. This will cause a similar number to be generated each time."

Ok, but why? This is yet another observation not quite an explanation. The answer may reside in the PRNG implementation.

Here because the program runs way below a second, three successive calls to `std::time ()` should return the same timestamp value so seed should (?) have the same value. In the other case, the global seed has the same value at each `randomReal ()` call.

There are many whys in my head, a lot of confusion too. For example, why reseeding is bad when reseeding should give even more randomness and so on. I will go deep in the mersenne implementation, if not now for sure in the near future. Thank you!



nascardriver

August 19, 2019 at 2:34 am · Reply.

A random number generator generates a predictable, seemingly random, sequence of numbers. The seed changes what this sequence looks like or where it starts. If you use the same seed twice, you'll get the same sequence twice.

> reseeding should give even more randomness

Random number generators can't be improved by changing the seed.



alfonso

August 19, 2019 at 11:47 pm · Reply.

Yes, I would expect that an unchangeable seed to give me similar results. But here in both cases the seed did not changed between successive `randomReal ()` calls. You can verify this using a code line:

```
1 | std::cout << "Seed: " << seed << '\n';
```

just before line 17 and you will see the same seed value for each `randomReal ()` call.

Though, when using global case it generates random different results. This I want to understand, why Alex MUST use a global variable. Of course, I can see the different output of programs, but I still do not understand why the global variable affects different the final output. Again, in both cases the seed does not change between successive `randomReal ()` calls. The reason must be somewhere else.

nascardriver



August 20, 2019 at 12:08 am · [Reply](#)

> the seed did not changed between successive randomReal () calls
That's the point. You're creating a new twister and seeding it with the same seed as last time every time you call `randomReal`. This will produce the same sequence and the first call to `realNumber` will give you the same number as it did last time.

With the global variable, the twister is only seeded once.

You might want to have a look at how simple pseudo random number generators are implemented (Don't look at the mersenne twister, unless you're up for some math and bits).



alfonso

August 20, 2019 at 11:19 pm · [Reply](#)

So it depends how the seed is used by mersenne twister. The seed is the same in all cases, but the seed is a sequence of many numbers and how that seed is used makes the difference. Seeding again it will use the same seed (sequence of numbers) in the same way. Seeding once it will use the same seed but in different ways between successive calls. So it looks like just scratching the surface and with your help. I can handle some math and bits but I can't handle class templates yet. Thank you!



nascardriver

August 21, 2019 at 4:58 am · [Reply](#)

> the seed is a sequence of many numbers
Not quite. The seed only initializes the generator and determines what the sequence will look like. Re-using a seed is like resetting the generator.

> I can handle some math and bits but I can't handle class templates yet.
You don't need to look at the cpp implementation. Mersenne twister is not language-specific.


https://en.wikipedia.org/wiki/Mersenne_Twister#Algorithmic_detail



Parsa

August 16, 2019 at 3:57 pm · [Reply](#)

How do you set the value of RAND_MAX since RANDMAX is a constant and is not modifiable?



Alex

August 16, 2019 at 9:19 pm · [Reply](#)

I think you just answered your own question. :)

You can't. If you need a random library that supports a wider range, use Mersenne Twister or something else.

