

## 7.16 — Lambda captures

BY NASCARDRIVER ON JANUARY 3RD, 2020 | LAST MODIFIED BY NASCARDRIVER ON JANUARY 26TH, 2020

### Capture clauses and capture by value

In the previous lesson, we introduced this example:

```
1  #include <algorithm>
2  #include <array>
3  #include <iostream>
4  #include <string_view>
5
6  int main()
7  {
8      std::array<std::string_view, 4> arr{ "apple", "banana", "walnut", "lemon" };
9
10     auto found{ std::find_if(arr.begin(), arr.end(),
11                             [](std::string_view str)
12                             {
13                                 return (str.find("nut") != std::string_view::npos);
14                             }) };
15
16     if (found == arr.end())
17     {
18         std::cout << "No nuts\n";
19     }
20     else
21     {
22         std::cout << "Found " << *found << '\n';
23     }
24
25     return 0;
26 }
```

Now, let's modify the nut example and let the user pick a substring to search for. This isn't as intuitive as you might expect.

```
1  #include <algorithm>
2  #include <array>
3  #include <iostream>
4  #include <string_view>
5  #include <string>
6
7  int main()
8  {
9      std::array<std::string_view, 4> arr{ "apple", "banana", "walnut", "lemon" };
10
11     // Ask the user what to search for.
12     std::cout << "search for: ";
13
14     std::string search{};
15     std::cin >> search;
16
17     auto found{ std::find_if(arr.begin(), arr.end(), [](std::string_view str) {
18         // Search for @search rather than "nut".
19     }) }
```

```

19     return (str.find(search) != std::string_view::npos); // Error: search not accessible in th
20 } };
```

```

21
22     if (found == arr.end())
23     {
24         std::cout << "Not found\n";
25     }
26     else
27     {
28         std::cout << "Found " << *found << '\n';
29     }
30
31     return 0;
32 }
```

This code won't compile. Unlike nested blocks, where any identifier defined in an outer block is accessible in the scope of the nested block, lambdas can only access specific kinds of identifiers: global identifiers, entities that are known at compile time, and entities with static storage duration. `search` fulfills none of these requirements, so the lambda can't see it. That's what the capture clause is there for.

## The capture clause

The **capture clause** is used to (indirectly) give a lambda access to variables available in the surrounding scope that it normally would not have access to. All we need to do is list the entities we want to access from within the lambda as part of the capture clause. In this case, we want to give our lambda access to the value of variable `search`, so we add it to the capture clause:

```

1  #include <algorithm>
2  #include <array>
3  #include <iostream>
4  #include <string_view>
5  #include <string>
6
7  int main()
8  {
9      std::array<std::string_view, 4> arr{ "apple", "banana", "walnut", "lemon" };
10
11      std::cout << "search for: ";
12
13      std::string search{};
14      std::cin >> search;
15
16      // Capture @search                                vvvvvv
17      auto found{ std::find_if(arr.begin(), arr.end(), [search](std::string_view str) {
18          return (str.find(search) != std::string_view::npos);
19      }) };
20
21      if (found == arr.end())
22      {
23          std::cout << "Not found\n";
24      }
25      else
26      {
27          std::cout << "Found " << *found << '\n';
28      }
29
30      return 0;
31 }
```

The user can now search for an element of our array.

## Output

```
search for: nana  
Found banana
```

## So how do captures actually work?

While it might look like our lambda in the example above is directly accessing the value of `main`'s `search` variable, this is not the case. Lambdas might look and function like nested blocks, but they work slightly differently (and the distinction is important).

When a lambda definition is executed, for each variable that the lambda captures, a clone of that variable is made (with an identical name) inside the lambda. These cloned variables are initialized from the outer scope variables of the same name at this point.

Thus, in the above example, when the lambda object is created, the lambda gets its own cloned variable named `search`. This cloned `search` has the same value as `main`'s `search`, so it behaves like we're accessing `main`'s `search`, but we're not.

While these cloned variable have the same name, they don't necessarily have the same type as the original variable. We'll explore this in the upcoming sections of this lesson.

### Key insight

The captured variables of a lambda are *clones* of the outer scope variables, not the actual variables.

### For advanced readers

Although lambdas look like functions, they're actually objects that can be called like functions (these are called **functors** -- we'll discuss how to create your own functors from scratch in a future lesson).

When the compiler encounters a lambda definition, it creates a custom object definition for the lambda. Each captured variable becomes a data member of the object.

At runtime, when the lambda definition is encountered, the lambda object is instantiated, and the members of the lambda are initialized at that point.

## Captures default to const value

By default, variables are captured by `const` value. This means when the lambda is created, the lambda captures a constant copy of the outer scope variable, which means that the lambda is not allowed to modify them. In the following example, we capture the variable `ammo` and try to decrement it.

```
1  #include <iostream>  
2  
3  int main()  
4  {  
5      int ammo{ 10 };  
6  
7      // Define a lambda and store it in a variable called "shoot".  
8      auto shoot{  
9          [ammo]() {
```

```
10     // Illegal, ammo was captured as a const copy.
11     --ammo;
12
13     std::cout << "Pew! " << ammo << " shot(s) left.\n";
14 }
15 };
16
17 // Call the lambda
18 shoot();
19
20 std::cout << ammo << " shot(s) left\n";
21
22 return 0;
23 }
```

In the above example, when we capture `ammo`, a new `const` variable with the same name and value is created in the lambda. We can't modify it, because it is `const`, which causes a compile error.

## Mutable capture by value

To allow modifications of variables that were captured by value, we can mark the lambda as `mutable`. The **`mutable`** keyword in this context removes the `const` qualification from *all* variables captured by value.

```
1  #include <iostream>
2
3  int main()
4  {
5      int ammo{ 10 };
6
7      auto shoot{
8          // Added mutable after the parameter list.
9          [ammo]() mutable {
10             // We're allowed to modify ammo now
11             --ammo;
12
13             std::cout << "Pew! " << ammo << " shot(s) left.\n";
14         }
15     };
16
17     shoot();
18     shoot();
19
20     std::cout << ammo << " shot(s) left\n";
21
22     return 0;
23 }
```

Output:

```
Pew! 9 shot(s) left.
Pew! 8 shot(s) left.
10 shot(s) left
```

While this now compiles, there's still a logic error. What happened? When the lambda was called, the lambda captured a *copy* of `ammo`. When the lambda decremented `ammo` from 10 to 9 to 8, it decremented its own copy, not the original value.

Note that the value of `ammo` is preserved across calls to the lambda!

## Capture by reference

Much like functions can change the value of arguments passed by reference, we can also capture variables by reference to allow our lambda to affect the value of the argument.

To capture a variable by reference, we prepend an ampersand (&) to the variable name in the capture. Unlike variables that are captured by value, variables that are captured by reference are non-const, unless the variable they're capturing is const. Capture by reference should be preferred over capture by value whenever you would normally prefer passing an argument to a function by reference (e.g. for non-fundamental types).

Here's the above code with ammo captured by reference:

```

1  #include <iostream>
2
3  int main()
4  {
5      int ammo{ 10 };
6
7      auto shoot{
8          // We don't need mutable anymore
9          [&ammo]() { // &ammo means ammo is captured by reference
10             // Changes to ammo will affect main's ammo
11             --ammo;
12
13             std::cout << "Pew! " << ammo << " shot(s) left.\n";
14         }
15     };
16
17     shoot();
18
19     std::cout << ammo << " shot(s) left\n";
20
21     return 0;
22 }
```

This produces the expected answer:

```
Pew! 9 shot(s) left.
9 shot(s) left
```

Now, let's use a reference capture to count how many comparisons `std::sort` makes when it sorts an array.

```

1  #include <algorithm>
2  #include <iostream>
3  #include <string>
4
5  struct Car
6  {
7      std::string make{};
8      std::string model{};
9  };
10
11 int main()
12 {
13     std::array<Car, 3> cars{ { { "Volkswagen", "Golf" },
14                               { "Toyota", "Corolla" },
15                               { "Honda", "Civic" } } };
16
17     int comparisons{ 0 };
18 }
```

```
19     std::sort(cars.begin(), cars.end(),
20             // Capture @comparisons by reference.
21             [&comparisons](const auto& a, const auto& b) {
22                 // We captured comparisons by reference. We can modify it without "mutable".
23                 ++comparisons;
24
25                 // Sort the cars by their make.
26                 return (a.make < b.make);
27             });
28
29     std::cout << "Comparisons: " << comparisons << '\n';
30
31     for (const auto& car : cars)
32     {
33         std::cout << car.make << ' ' << car.model << '\n';
34     }
35
36     return 0;
37 }
```

Output

```
Comparisons: 2
Honda Civic
Toyota Corolla
Volkswagen Golf
```

---

## Capturing multiple variables

Multiple variables can be captured by separating them with a comma. This can include a mix of variables captured by value or by reference:

```
1  int health{ 33 };
2  int armor{ 100 };
3  std::vector<CEntity> enemies{};
4
5  // Capture health and armor by value, and enemies by reference.
6  [health, armor, &enemies](){};
```

---

## Default captures

Having to explicitly list the variables you want to capture can be burdensome. If you modify your lambda, you may forget to add or remove captured variables. Fortunately, we can enlist the compiler's help to auto-generate a list of variables we need to capture.

A **default capture** (also called a **capture-default**) captures all variables that are mentioned in the lambda. Variables not mentioned in the lambda are not captured if a default capture is used.

To capture all used variables by value, use a capture value of `=`.

To capture all used variables by reference, use a capture value of `&`.

Here's an example of using a default capture by value:

```
1  #include <array>
2  #include <iostream>
3
4  int main()
5  {
```

```

6   std::array areas{ 100, 25, 121, 40, 56 };
7
8   int width{};
9   int height{};
10
11  std::cout << "Enter width and height: ";
12  std::cin >> width >> height;
13
14  auto found{ std::find_if(areas.begin(), areas.end(),
15                        [=](int knownArea) { // will default capture width and height by va
16                        return (width * height == knownArea); // because they're mentione
17                        }) };
18
19  if (found == areas.end())
20  {
21      std::cout << "I don't know this area :(\n";
22  }
23  else
24  {
25      std::cout << "Area found :)\n";
26  }
27
28  return 0;
29 }

```

Default captures can be mixed with normal captures. We can capture some variables by value and others by reference, but each variable can only be captured once.

```

1   int health{ 33 };
2   int armor{ 100 };
3   std::vector<CEntity> enemies{};
4
5   // Capture health and armor by value, and enemies by reference.
6   [health, armor, &enemies](){};
7
8   // Capture enemies by reference and everything else by value.
9   [=, &enemies](){};
10
11  // Capture armor by value and everything else by reference.
12  [&, armor](){};
13
14  // Illegal, we already said we want to capture everything by reference.
15  [&, &armor](){};
16
17  // Illegal, we already said we want to capture everything by value.
18  [=, armor](){};
19
20  // Illegal, armor appears twice.
21  [armor, &health, &armor](){};
22
23  // Illegal, the default capture has to be the first element in the capture group.
24  [armor, &](){};

```

## Defining new variables in the lambda-capture

Sometimes we want to capture a variable with a slight modification or declare a new variable that is only visible in the scope of the lambda. We can do so by defining a variable in the lambda-capture without specifying its type.

```

1   #include <array>
2   #include <iostream>
3

```

```

4  int main()
5  {
6      std::array areas{ 100, 25, 121, 40, 56 };
7
8      int width{};
9      int height{};
10
11     std::cout << "Enter width and height: ";
12     std::cin >> width >> height;
13
14     // We store areas, but the user entered width and height.
15     // We need to calculate the area before we can search for it.
16     auto found{ std::find_if(areas.begin(), areas.end(),
17                             // Declare a new variable that's visible only to the lambda.
18                             // The type of userArea is automatically deduced to int.
19                             [userArea{ width * height }](int knownArea) {
20                                 return (userArea == knownArea);
21                             }) };
22
23     if (found == areas.end())
24     {
25         std::cout << "I don't know this area :(\n";
26     }
27     else
28     {
29         std::cout << "Area found :)\n";
30     }
31
32     return 0;
33 }

```

`userArea` will only be calculated once when the lambda is defined. The calculated area is stored in the lambda object and is the same for every call. If a lambda is mutable and modifies a variable that was defined in the capture, the original value will be overridden.

### Best practice

Only initialize variables in the capture if their value is short and their type is obvious. Otherwise it's best to define the variable outside of the lambda and capture it.

## Dangling captured variables

Variables are captured at the point where the lambda is defined. If a variable captured by reference dies before the lambda, the lambda will be left holding a dangling reference.

For example:

```

1  #include <iostream>
2  #include <string>
3
4  // returns a lambda
5  auto makeWalrus(const std::string& name)
6  {
7      // Capture name by reference and return the lambda.
8      return [&]() {
9          std::cout << "I am a walrus, my name is " << name << '\n'; // Undefined behavior
10     };
11 }

```



```
12
13 int main()
14 {
15     // Create a new walrus whose name is Roofus.
16     // sayName is the lambda returned by makeWalrus.
17     auto sayName{ makeWalrus("Roofus") };
18
19     // Call the lambda function that makeWalrus returned.
20     sayName();
21
22     return 0;
23 }
```

The call to `makeWalrus` creates a temporary `std::string` from the string literal “Roofus”. The lambda in `makeWalrus` captures the temporary string by reference. The temporary string dies when `makeWalrus` returns, but the lambda still references it. Then when we call `sayName`, the dangling reference is accessed, causing undefined behavior.

Note that this also happens if `name` is passed to `makeWalrus` by value. The variable `name` still dies at the end of `makeWalrus`, and the lambda is left holding a dangling reference.

### Warning

Be extra careful when you capture variables by reference, especially with a default reference capture. The captured variables must outlive the lambda.

If we want the captured name to be valid when the lambda is used, we need to capture it by value instead (either explicitly or using a default-capture by value).

## Unintended copies of mutable lambdas

Because lambdas are objects, they can be copied. In some cases, this can cause problems. Consider the following code:

```
1  #include <iostream>
2
3  int main()
4  {
5      int i{ 0 };
6
7      // Create a new lambda named count
8      auto count{ [i]() mutable {
9          std::cout << ++i << '\n';
10     } };
11
12     count(); // invoke count
13
14     auto otherCount{ count }; // create a copy of count
15
16     // invoke both count and the copy
17     count();
18     otherCount();
19
20     return 0;
21 }
```

Output

1  
2  
2

Rather than printing 1, 2, 3, the code prints 2 twice. When we created `otherCount` as a copy of `count`, we created a copy of `count` in its current state. `count`'s `i` was 1, so `otherCount`'s `i` is 1 as well. Since `otherCount` is a copy of `count`, they each have their own `i`.

Now let's take a look at a slightly less obvious example:

```

1  #include <iostream>
2  #include <functional>
3
4  void invoke(const std::function<void(void)>& fn)
5  {
6      fn();
7  }
8
9  int main()
10 {
11     int i{ 0 };
12
13     // Increments and prints its local copy of @i.
14     auto count{ [i]() mutable {
15         std::cout << ++i << '\n';
16     } };
17
18     invoke(count);
19     invoke(count);
20     invoke(count);
21
22     return 0;
23 }
```

Output:

1  
1  
1

This exhibits the same problem as the prior example in a more obscure form. When `std::function` is created with a lambda, the `std::function` internally makes a copy of the lambda object. Thus, our call to `fn()` is actually being executed on the copy of our lambda, not the actual lambda.

If we need to pass a mutable lambda, and want to avoid the possibility of inadvertent copies being made, there are two options. One option is to use a non-capturing lambda instead -- in the above case, we could remove the capture and track our state using a static local variable instead. But static local variables can be difficult to keep track of and make our code less readable. A better option is to prevent copies of our lambda from being made in the first place. But since we can't affect how `std::function` (or other standard library functions or objects) are implemented, how can we do this?

Fortunately, C++ provides a convenient type (as part of the `<functional>` header) called `std::ref` that allows us to pass a normal type as if it were a reference. By wrapping our lambda in a `std::ref`, whenever anybody tries to make a copy of our lambda, they'll make a copy of the reference instead, which will copy the reference rather than the actual object.

Here's our updated code using `std::ref`:

```
1  #include <iostream>
2  #include <functional>
3
4  void invoke(const std::function<void(void)> &fn)
5  {
6      fn();
7  }
8
9  int main()
10 {
11     int i{ 0 };
12
13     // Increments and prints its local copy of @i.
14     auto count{ [i]() mutable {
15         std::cout << ++i << '\n';
16     } };
17
18     // std::ref(count) ensures count is treated like a reference
19     // thus, anything that tries to copy count will actually copy the reference
20     // ensuring that only one count exists
21     invoke(std::ref(count));
22     invoke(std::ref(count));
23     invoke(std::ref(count));
24
25     return 0;
26 }
```

Our output is now as expected:

```
1
2
3
```

Note that the output doesn't change even if `invoke` takes `fn` by value. `std::function` doesn't create a copy of the lambda if we create it with `std::ref`.

### Rule

Standard library functions may copy function objects (reminder: lambdas are function objects). If you want to provide lambdas with mutable captured variables, pass them by reference using `std::ref`.

### Best practice

Try to avoid lambdas with states altogether. Stateless lambdas are easier to understand and don't suffer from the above issues, as well as more dangerous issues that arise when you add parallel execution.

## Quiz time

### Question #1

Which of the following variables can be used by the lambda in `main` without explicitly capturing them?

```
1  int i{};
2  static int j{};
3
4  int getValue()
5  {
6      return 0;
7  }
8
9  int main()
10 {
11     int a{};
12     constexpr int b{};
13     static int c{};
14     static constexpr int d{};
15     const int e{};
16     const int f{ getValue() };
17     static const int g{};
18     static const int h{ getValue() };
19
20     [](){
21         // Try to use the variables without explicitly capturing them.
22         a;
23         b;
24         c;
25         d;
26         e;
27         f;
28         g;
29         h;
30         i;
31         j;
32     }();
33
34     return 0;
35 }
```

### Show Solution

### Question #2

What does the following code print? Don't run the code, work it out in your head.

```
1  #include <iostream>
2  #include <string>
3
4  int main()
5  {
6      std::string favoriteFruit{ "grapes" };
7
8      auto printFavoriteFruit{
9          [=]() {
10             std::cout << "I like " << favoriteFruit << '\n';
11         }
12     };
13
14     favoriteFruit = "bananas with chocolate";
15
16     printFavoriteFruit();
17
18     return 0;
19 }
```

**Show Solution****Question #3**

We're going to write a little game with square numbers (numbers which can be created by multiplying an integer with itself (1, 2, 4, 9, 16, 25, ...)).

Ask the user to input 2 numbers, the first is the square root of the number to start at, the second in the amount to numbers to generate. Generate a random integer from 2 to 4, and square numbers in the range that was chosen by the user. Multiply each square number by the random number. You can assume that the user enters valid numbers.

The user has to calculate which numbers have been generated. The program checks if the user guessed correctly and removes the guessed number from the list. If the user guessed wrong, the game is over and the program prints the number that was closest to the user's final guess, but only if the final guess was not off by more than 4.

Here are a couple of sample sessions to give you a better understanding of how the game works:

```
Start where? 4
How many? 8
I generated 8 square numbers. Do you know what each number is after multiplying it by 2?
> 32
Nice! 7 number(s) left.
> 72
Nice! 6 number(s) left.
> 50
Nice! 5 number(s) left.
> 126
126 is wrong! Try 128 next time.
```

- The user chose to start at 4 and wants to play with 8 numbers.
- Each square number will be multiplied by 2. 2 was randomly chosen by the program.
- The program generates 8 square number, starting with 4 as a base:
- 16 25 36 49 64 81 100 121
- But each number is multiplied by 2, so we get:
- 32 50 72 98 128 162 200 242
- Now the user starts to guess. The order in which in guesses are entered doesn't matter.
- 32 is in the list.
- 72 is in the list.
- 126 is not in the list, the user loses. There is a number in the list (128) that is not more than 4 away from the user's guess, so that number is printed.

```
Start where? 1
How many? 3
I generated 3 square numbers. Do you know what each number is after multiplying it by 4?
> 4
Nice! 2 numbers left.
> 16
Nice! 1 numbers left.
> 36
Nice! You found all numbers, good job!
```

- The user chose to start at 1 and wants to play with 3 numbers.

- Each square number will be multiplied by 4.
- The program generates these square numbers:
- 1 4 9
- Multiplied by 4
- 4 16 36
- The user guesses all numbers correctly and wins the game.

Start where? 2

How many? 2

I generated 2 square numbers. Do you know what each number is after multiplying it by 4?

> 21

21 is wrong!

- The user chose to start at 2 and wants to play with 2 numbers.
- Each square number will be multiplied by 4.
- The program generates these numbers:
- 16 36
- The user guesses 21 and loses. 21 is not close enough to any of the remaining numbers, so no number is printed.

Use `std::find` (**7.8 -- Function Pointers**) to search for a number in the list.

Use `std::vector::erase` to remove an element, e.g.

```
1 auto found{ std::find(/* ... */) };
2
3 // Make sure the element was found
4
5 myVector.erase(found);
```

Use `std::min_element` and a lambda to find the number closest to the user's guess. `std::min_element` works analogous to `std::max_element` from the previous quiz.

**Show Hint**

**Show Solution**



**7.x -- Chapter 7 comprehensive quiz**



**Index**



**7.15 -- Introduction to lambdas (anonymous functions)**

 [C++ TUTORIAL](#) |  [PRINT THIS POST](#)

## 11 comments to 7.16 — Lambda captures



Suyash

[January 25, 2020 at 1:04 pm](#) · [Reply](#)

Here's my solution to the problem... Will patiently wait for your review...

Is this the proper usage of `static` keyword in the context of functions if I wish to prevent it from being used in other files of the project. Or, is it redundant due to them not being declared in the header? And, is there any benefit in terms of performance (either runtime or compile/build time) when using `static` (w/ functions)?

Another question is why can't a person with a pre-C++17 compiler do this exercise?

### 1. guess\_squares.h

```

1  #ifndef GUESS_SQUARES_H
2  #define GUESS_SQUARES_H
3
4  namespace guess_square
5  {
6      /*
7       * Simulate a round of "Guess Squares" game.
8       */
9      void playGuessSquares();
10 }
11
12 #endif // !GUESS_SQUARES_H

```

### 2. guess\_squares.cpp

```

1  #include "guess_squares.h"
2
3  #include <algorithm>      // std::find() & std::min_element()
4  #include <cmath>          // std::pow() & std::abs()
5  #include <ctime>          // std::time
6  #include <iostream>       // std::cout & std::cin
7  #include <limits>         // std::numeric_limits<std::streamsize>::max()
8  #include <random>         // std::mt19937 & std::uniform_int_distribution
9  #include <string_view>    // std::string_view
10 #include <vector>         // std::vector
11

```

```
12 namespace guess_square
13 {
14     /*
15     Get the user to enter a positive integer.
16     */
17     static int getUserInput(const std::string_view query_str)
18     {
19         constexpr int origin{ 0 };
20
21         do
22         {
23             std::cout << query_str << ' ';
24             int n{};
25             std::cin >> n;
26
27             if (std::cin.fail())
28             {
29                 std::cin.clear();
30                 std::cin.ignore(std::numeric_limits<std::streamsize>::max(), '\n');
31                 std::cout << "Invalid Type. Enter a positive integer.\n";
32             }
33             else if (n <= origin)
34             {
35                 std::cout << "Invalid Value. Enter a positive integer.\n";
36             }
37             else
38             {
39                 return n;
40             }
41         } while (true);
42     }
43
44     /*
45     Generate a random integer between 2 and 4.
46     */
47     static int getRandomMultiplyingFactor()
48     {
49         constexpr int min{ 2 };
50         constexpr int max{ 4 };
51
52         static std::mt19937 mersenne{ static_cast<std::mt19937::result_type>(std::time(nullptr)) };
53         std::uniform_int_distribution range{ min, max };
54         return range(mersenne);
55     }
56
57     /*
58     Return a vector array containing a list of numbers generated by
59     multiplying the square of a number (present in the range) with a
60     factor.
61     */
62     static std::vector<int> generateSquaresList(const int start, const int size, const int
63     {
64         std::vector<int> vec{};
65
66         vec.resize(size);
67
68         int element{ start };
69
70         for (int index{}; index < vec.size(); ++index)
71         {
72             vec[index] = std::pow(element, 2) * factor;
73             ++element;
```



```

74     }
75
76     return vec;
77 }
78
79 void playGuessSquares()
80 {
81     // Ask the user to enter the square root of the number to start at.
82     int start_sqrt{ getUserInput("Start where?") };
83
84     // Ask the user to enter the amount of numbers to generate (range).
85     int range_size{ getUserInput("How many?") };
86
87     const int multiplying_factor{ getRandomMultiplyingFactor() };
88
89     // Generate the square numbers in the range that was chosen by the user.
90     std::vector squares_list{ generateSquaresList(start_sqrt, range_size, multiplying_
91
92     std::cout << "I generated " << range_size << " square numbers. Do you know what"
93         << " each number is after multiplying it by " << multiplying_factor << "?\n";
94
95     do
96     {
97         // Let the user guess a particular number from the list of generated numbers.
98         int guessed_square{ getUserInput(">") };
99         std::vector<int>::iterator list_iter{ std::find(squares_list.begin(), squares_
100
101         // If the user guesses right, remove that number from the list.
102         if (list_iter != squares_list.end())
103         {
104             squares_list.erase(list_iter);
105
106             // First, decrement `range_size` by 1 & then, perform the comparison.
107             if (--range_size)
108                 std::cout << "Nice! " << range_size << " number(s) left.\n";
109         }
110         // Otherwise, if the user guessed wrong, the game is over
111         // and the program prints the number that was closest to
112         // the user's final guess, but only if the final guess was
113         // not off by more than 4.
114         else
115         {
116             std::cout << guessed_square << " is wrong.";
117
118             auto closest_number_to_guess{ *std::min_element(squares_list.begin(), squa
119                 [guessed_square](const auto a, const auto b)
120                 {
121                     int distance_to_a{ std::abs(a - guessed_square)};
122                     int distance_to_b{ std::abs(b - guessed_square)};
123
124                     return distance_to_a < distance_to_b;
125                 }
126             } };
127
128             constexpr int epsilon{ 4 };
129             if (std::abs(closest_number_to_guess - guessed_square) <= epsilon)
130             {
131                 std::cout << " Try " << closest_number_to_guess << " next time.\n";
132                 return;
133             }
134             else
135             {

```

```

136         std::cout << '\n';
137         return;
138     }
139 }
140 } while (range_size);
141
142 std::cout << "Nice! You found all numbers, good job!\n";
143 }
144 }
```

### 3. main.cpp

```

1  #include "guess_squares.h"
2
3  int main()
4  {
5      guess_square::playGuessSquares();
6
7      return 0;
8  }
```

#### Side notes:

1) Yesterday, I observed that the notification emails were considered as 'spam' and hence, were not showing up in my inbox... Solved the issue.

2) Possible bug/issue - I also received an email notifying that @nascardriver has replied to my comment on the previous chapter (7.15 - Anonymous Functions) but whenever I visit the page (via the 'View comment' link), I don't see his reply. Just my post.



nascardriver

January 26, 2020 at 1:33 am · Reply

> Another question is why can't a person with a pre-C++17 compiler do this exercise? The quiz used to use `std::generate`, which is a C++17 function. It has been removed from the quiz, but the comment stayed. I removed it, thanks for pointing it out! (The solution still uses C++17 features, but can easily be ported to C++14).

Your use of `static` is correct. There's no noticeable performance difference resulting from the use of `static` functions.

- Nice sorted includes :)
- I like the query string of your `getUserInput` function.
- Line 64-66: You can create a `std::vector` with a specific size from the get go. `std::vector vec(size)`. Note the use of direct initialization here. Brace initialization would add a single element with value `size`.
- Line 70+: You don't need an index. A range-based for-loop is easier to read.
- Line 99: It doesn't matter what kind of iterator this is. `auto` is fine.
- `range\_size` mirrors `squares\_list.size()`.

Very good code overall!

2)

Fixed, my comment was pending approval. Either I unapprove my own comments or I'm triggering a filter. I'll have to look into that. You should now be able to see the comment.



Suyash

January 24, 2020 at 12:38 am · Reply

Okay, I still need to attempt the exercises but I have a question... Okay, the concept of lambda captures seem like a new one... But, it seems overly complicated approach to achieve a simple task of getting access to the variables present in the enclosing scope...

For example, let's again look at lambda capture used in the "search nut" example...

```

1  #include <algorithm>
2  #include <array>
3  #include <iostream>
4  #include <string_view>
5  #include <string>
6
7  int main()
8  {
9      std::array<std::string_view, 4> arr{ "apple", "banana", "walnut", "lemon" };
10
11     std::cout << "search for: ";
12
13     std::string search{};
14     std::cin >> search;
15
16     /*
17     Lambda Functions definitions here
18     */
19     if (found == arr.end())
20     {
21         std::cout << "Not found\n";
22     }
23     else
24     {
25         std::cout << "Found " << *found << '\n';
26     }
27
28     return 0;
29 }

1  // Lambda 1: Capture @search using capture          vvvvvv
2  auto found{ std::find_if(arr.begin(), arr.end(), [search](std::string_view str) {
3      return (str.find(search) != std::string_view::npos);
4  }} };

```

I am just curious about the reason being the concept of lambda captures... For a person coming from another language where the same concept is implemented in a... more sophisticated fashion, this just feels a bit unintuitive...

Why the above code is preferred over something more simple like the following? Is it due to lambdas being functors (function objects)? I know that my following code snippet doesn't work but that raises another question in my head...

Is it not possible for us to call/invoke lambda functions like regular functions? And, should lambdas be only used as (a function) parameter to another function?

Probably, it's due to the fact that functions can't be nested and since functors already use the overloaded () operator to declare parameters & thus, they can't be used to make calls to them...

```

1  // Lambda 2: Capture @search as another parameter
2  auto found{ std::find_if(arr.begin(), arr.end(), [](std::string_view str, std::string sear
3      return (str.find(search) != std::string_view::npos);
4  }} };

```

I know I sound like a mess but I am just confused with the idea of lambda captures... I will patiently wait for your answer...



nascardriver

[January 24, 2020 at 2:31 am · Reply](#)

Your concerns are valid and your suggested code isn't unpopular either.

> it seems overly complicated approach to achieve a simple task of getting access to the variables  
C++ is a language that's close to the system. You can do a lot, and you have a lot of control over what's happening on a low-level. But to achieve that level of control, things have to be complicated at times. Lambdas are objects, blindly copying variables into them to make them accessible would be counter-intuitive, so we get to choose between what we want to capture and how to capture it. This comes at the cost of a more verbose syntax than other languages have.

> Why the above code is preferred over something more simple like the following?

If you want your lambda to take another parameter, then that parameter has to get its value from somewhere. `std::find_if` passes a single container element to the function, because that's all the function should need. If the function needs more, `std::find_if` can't know that, and you have to get the value from somewhere else, enter captures.

If you called the lambda yourself, you could add another parameter rather than capturing a variable.

> Is it not possible for us to call/invoke lambda functions like regular functions?

Regular functions and lambdas are used exactly the same from a syntax point of view. Internally, they're different, but that doesn't usually matter.

> should lambdas be only used as (a function) parameter to another function?

You can use them wherever you want (Almost). Most of the time, you'll need them as an argument for a callback or similar.

> functors already use the overloaded () operator to declare parameters & thus, they can't be used to make calls to them

I'm not sure I understand what you're saying. Lambdas can be called directly, that's done for complex initializations for example.

```
1  std::string str{
2      []() { // The parameter list can be omitted. I kept it.
3          std::string str{};
4          std::cin >> str;
5          return str;
6      }(); // Call immediately
7  };
```



Suyash

[January 24, 2020 at 2:57 am · Reply](#)

Thanks for the quick response, @nascardriver... Your answer cleared a lot of doubts regarding the concept of Lambdas & its implementation in the language...

C++ could be considered as my third language (at the current moment after Python/JavaScript) and since I have to constantly switch between them means that I end up mixing a lot of stuff (especially in the case of C++)...

While I love the fact that C++ has been able to introduce modern programming concepts in their core language, the verbose syntax is something that I have come to dislike (due to spending so much time in Python) but at the same time, I guess, I am also learning it to appreciate it on how much easier, it has become for me to write self-documenting code and how I can use it to describe my intent...

Every language has its pros & cons... But, even though I have only spent a short time in C++, I do like how it has helped me become a more aware programmer... While I don't always like the motto

of "Trust the programmer!" especially because I have so easily been burnt by it (on my previous attempts using C++), but now I am glad for it...

By the way, how can I quote a post (or if possible, a particular paragraph) written by someone?



nascardriver

[January 24, 2020 at 3:19 am · Reply](#)

Apart from code tags, there's nothing special about comments (Authors can use html, that's why we can use links). If I want to quote something, I copy it and add

> in front.



Fan

[January 10, 2020 at 11:38 am · Reply](#)

In the last example, is it possible to avoid copying of count by changing the type of the argument of invoke to auto&?

Also, in the section "Dangling captured variables", there is an extra period after the phrase "dangling reference".



nascardriver

[January 12, 2020 at 3:09 am · Reply](#)

> is it possible to avoid copying of count by changing the type of the argument of invoke to auto&?

Yes, that would work. Keep in mind that regular functions can't have `auto` parameters (Until C++20). Though, then the user of `invoke` doesn't know what kind of argument they can pass to `invoke`. Furthermore, that solution only works if we can modify `invoke`. If it's a function in a library, we can't modify its parameters.

> there is an extra period

Fixed, thanks!



Wilibert

[January 8, 2020 at 6:35 pm · Reply](#)

Hi, I have a question regarding the use of `std::min_element`... I was getting an error when I was trying this for quiz 3:

```
1 | int* found{ std::min_element(myVector.begin(), myVector.end(),           //myVector holds the nu
2 | [guess](const int a, const int b) -> bool {
3 |     return (std::abs(a - guess) < std::abs(b - guess));
4 | })
5 | };
```

Neither making found an int nor an int pointer worked. After looking at the solution, I noticed you were using a dereference operator right before `std::min_element` like so:

```
1 | int found{ *std::min_element(myVector.begin(), myVector.end(),
2 | [guess](const int a, const int b) -> bool {
3 |     return (std::abs(a - guess) < std::abs(b - guess));
4 | })
5 | };
```

And now I'm confused. Why would that work? Is it because the `*` operator for `std::min_element` calls a specially overloaded function that turns whatever type `std::min_element` normally returns into an `int`? Otherwise I can't see why an ordinary `int` pointer wouldn't work. I also tried changing `found`'s type to `auto` and dereferencing it whenever I needed it later and it also worked. So... normal `int` pointer cannot contain `std::min_element`'s return type but you can dereference that type anywhere and the compiler will do the right thing? Oh and thanks a lot for the lesson, I still remember when my friends were complaining after learning about lambdas (they have a proper college computer science education and I don't) so it means a lot to me to be able to learn this.



nascardriver

January 9, 2020 at 4:10 am · Reply

`std::min_element` returns an iterator. If you missed the lesson on iterators (added 17. dec 2019), it's [here](#).

Interacting with iterators is very similar to interacting with pointers, hence `*operator` works. But that `*operator` isn't a pointer indirection, it calls a special function of the iterator, which returns the element it's currently at.

When you use `*std::min_element(...)`, you're getting the element right away (If it exists, an error otherwise). When you use `auto`, you're storing the iterator for later and can use `*operator` on the iterator. There is no conversion from the iterator to a pointer.



Wilibert

January 9, 2020 at 11:23 am · Reply

Okay, that makes sense. It seems like I was confusing iterators as some form of pointer since they work similarly. I'll review the lesson on iterators to make sure I understand.

Thanks for the help!